

# Introduction to SQL

The background of the slide is an abstract composition. It features a light gray grid of thin, curved lines that sweep across the frame. Overlaid on this are several thick, vibrant lines in shades of magenta, lime green, and dark purple. These lines are dynamic, curving and intersecting. Additionally, there are thinner orange lines with small square markers, resembling a data series or a path. The overall effect is one of modern, technological energy.

# SQL (structured query language)

## ⚡ History

1974: first paper by Chamberlin&Boyce

SQL 92 (SQL 2): joins, outer-joins, ...

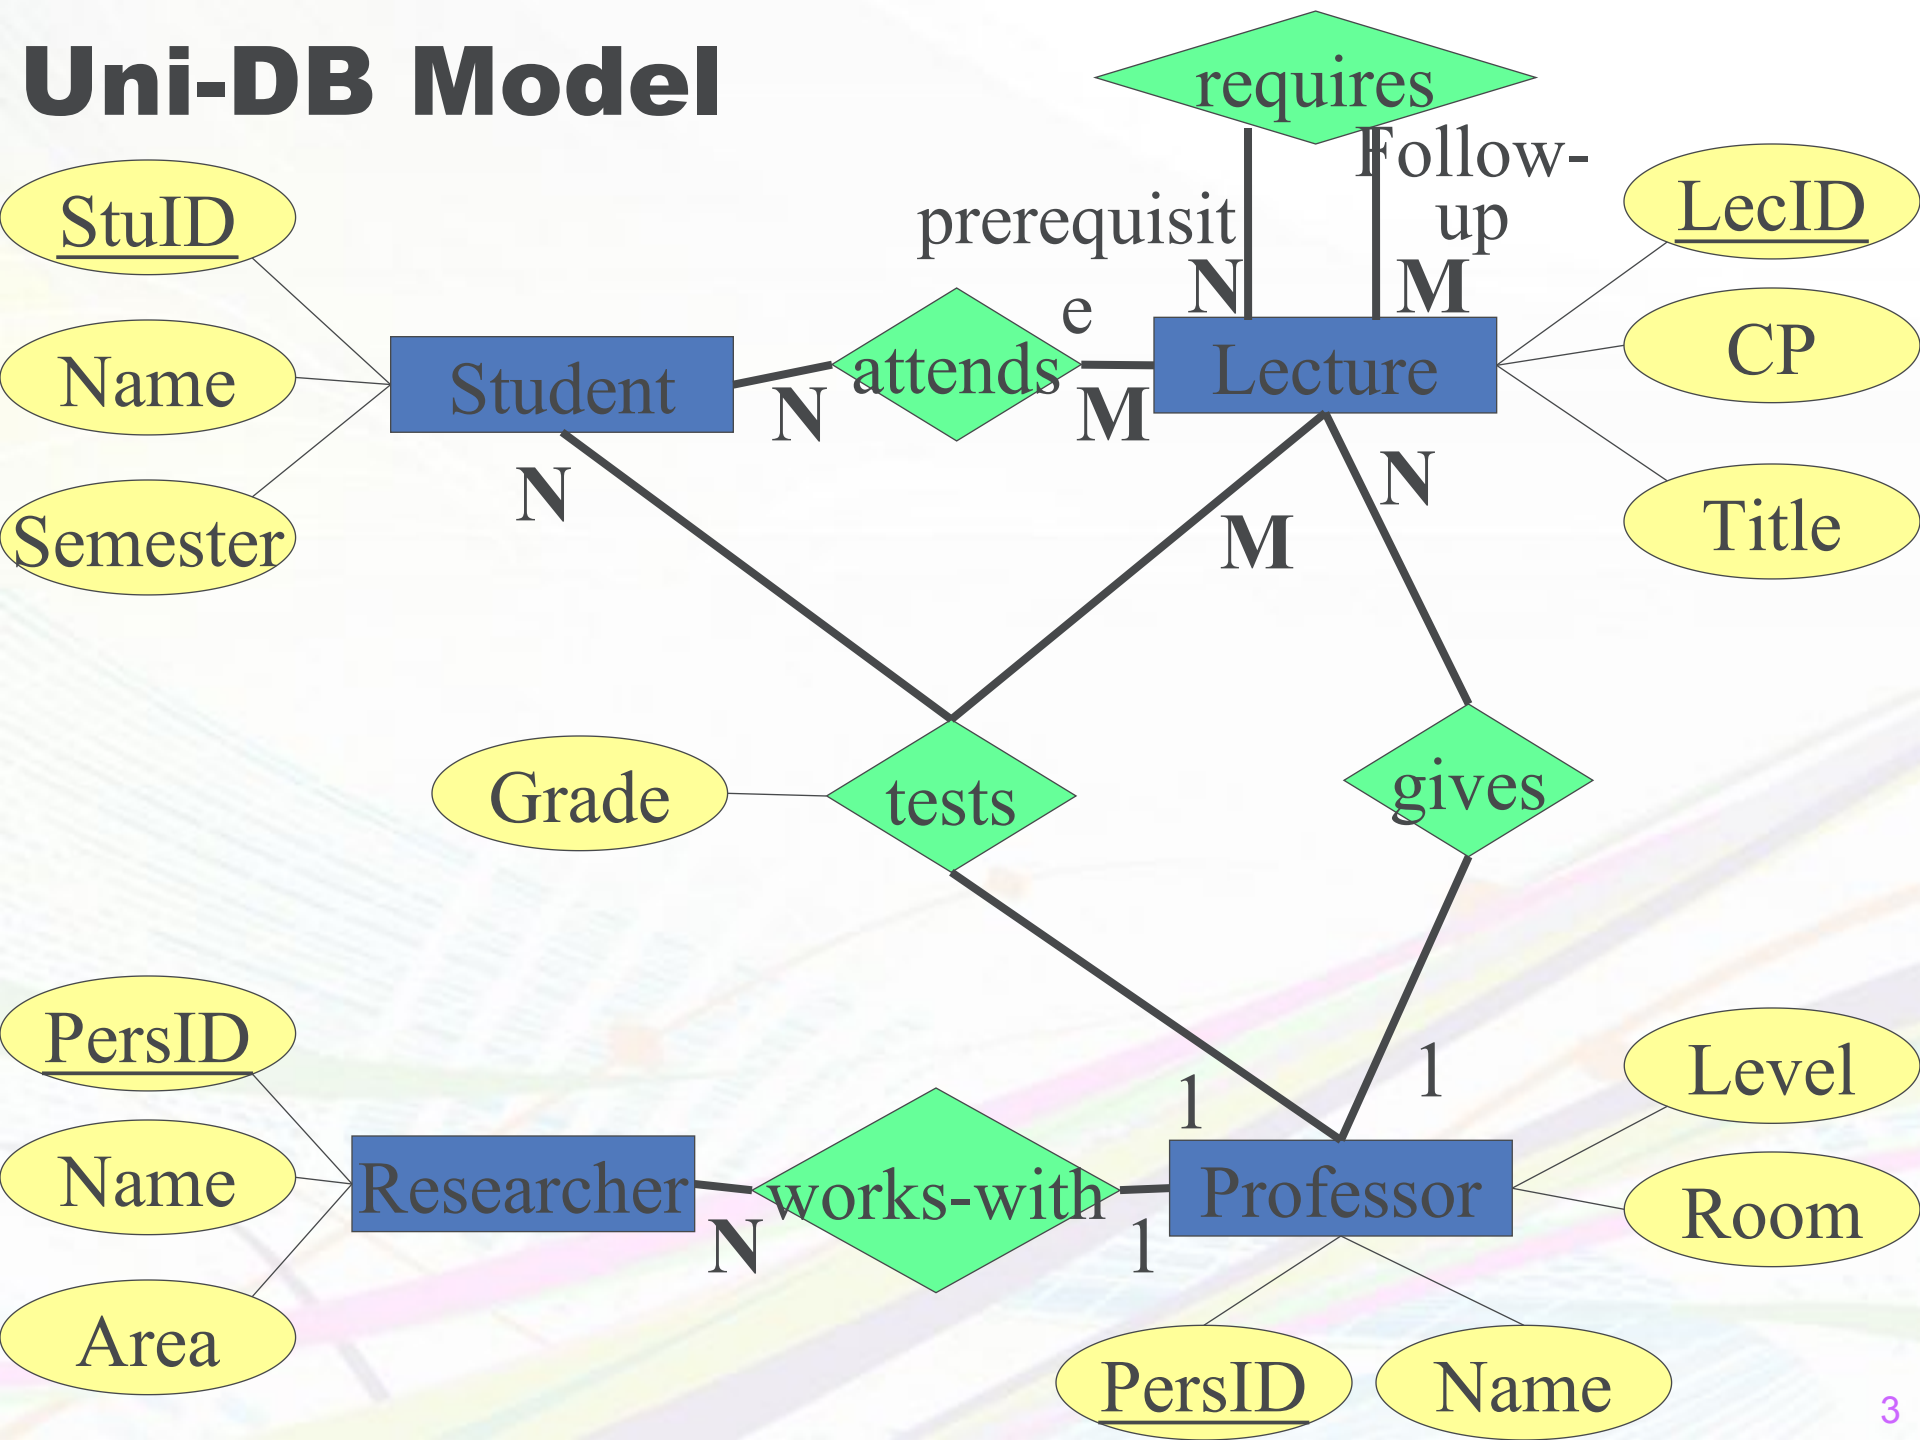
SQL 3: object-relational extensions

SQL/XML, etc.: domain-specific extensions

## ⚡ A family of standards

- Data definition language (DDL) - schemas
- Data manipulation language (DML) - updates
- Query language (Query) - reads

# Uni-DB Model



# (Simple) Data Definition with SQL

## Data types

- ↙ **character** (*n*), **char** (*n*)
- ↙ **character varying** (*n*), **varchar** (*n*)
- ↙ **numeric** (*p,s*), **integer**
- ↙ **blob** or **raw** for large binaries
- ↙ **clob** for large string values
- ↙ **date**

### Create Tables

```
create table Professor  
(PersID integer not null,  
  Name varchar (30) not null  
  Level character (2) default 'AP');
```



# DDL (ctd.)

## Delete a Table

↙ **drop table** Professor;

## Modify the structure of a Table

↙ **alter table** Professor **add column**(age integer);

## Management of Indexes (Performance tuning)

↙ **create index** myIndex **on** Professor(name, age);

↙ **drop index** myIndex;

# Updates (DML)

## Insert Tuples

**insert into** attends

**select** StuID, LecID  
**from** Student, Lecture  
**where** Title= 'DMD' ;

**insert into** Student (StuID, Name)  
**values** (28121, `Archimedes`);

Student		
StuID	Name	Semester
⋮	⋮	⋮
29120	Theophrastos	2
29555	Feuerbach	2
28121	Archimedes	-



Null

## Sequence Types (Automatic Increment for Surrogates)

↙ create sequence PersID\_seq increment by 1 start with 1;

↙ insert into Professor(PersID, Name)  
values(*PersNr\_seq.nextval*, 'Roscoe');

↙ Syntax is vendor dependent

E. g. , AUTO-INCREMENT Option in MySQL

Syntax above was standardized in SQL 2003

*what about PostgreSQL??*



# Updates (ctd.)

## Delete tuples

```
delete Student  
where Semester > 13;
```

## Update tuples

```
update Student  
set Semester= Semester + 1;
```

# Queries

```
select    PersID, Name
from      Professor
where     Level = 'FP';
```

PersID	Name
2125	Qiang
2126	Sadegh
2136	Joo
2137	Someone

# Queries: Sorting

```
select PersID, Name, Level  
from Professor  
order by Level desc, Name asc;
```

PersID	Name	Level
2136	Curie	FP
2137	sadegh	FP
2126	someone	FP
2127	Kopernikus	AP
2133	Popper	AP

# Duplicate Elimination

**select distinct Level**  
**from Professor**

Level
AP
FP

# Queries: Joins

Who teaches DMD?

**select** Name

**from** Professor, Lecture

**where** PersID = ProfID **and** Title = 'DMD' ;

$\Pi_{\text{Name}}(\sigma_{\text{PersID} = \text{ProfID} \wedge \text{Title} = \text{'DMD'}}(\text{Professor} \times \text{Lecture}))$

N.B.: Renamed Lecture.PersID to ProfID.

Will show later how this can be done as part of a query.



# Joins

Professor			
PersID	Name	Level	Room
2125	Qiang	AP	226
2126	Sadegh	FP	232
⋮	⋮	⋮	⋮
2137	someother	FP	7

Lecture			
LecID	Title	CP	ProfID
5001	DSP	4	2137
5041	OOP	4	2125
⋮	⋮	⋮	⋮
5049	DMD	2	2125
⋮	⋮	⋮	⋮
4630	whatever	4	2137

X



PerID	Name	Level	Room	LecID	Title	CP	ProfID
2125	Qiang	AP	226	5001	DSP	4	2137
1225	Manuel	FP	226	5041	OOP	4	2125
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2125	Qiang	FP	226	5049	DMD	2	2125
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2126	Sadegh	FP	232	5001	ML	4	2137
2126	Sadegh	FP	232	5041	IR	4	2125
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2137	someother	FP	7	4630	whatever	4	2137

↓  $\sigma$

PersID	Name	Level	Room	LecID	Title	CP	ProfID
2125	Qiang	AP	226	5049	DMD	2	2125

↓  $\pi$

Name
Qiang

# SQL -> Relational Algebra


SQL

Relational Algebra

**select**  $A_1, \dots, A_n$   
**from**  $R_1, \dots, R_k$   
**where**  $P$ ;



$$\Pi_{A_1, \dots, A_n}(\sigma_P(R_1 \times \dots \times R_k))$$



# Joins and Tuple Variables

Who attends which lecture?

```
select Name, Title
from Student, attends, Lecture
where Student.StuID = attends.StuID and
        attends.LecID = Lecture.LecID;
```

**Alternative:**

```
select s.Name, l.Title
from Student s, attends a, Lecture l
where s.StuID= a.StuID and
        a.LecID = l.LecID;
```

# Rename of Attributes

Give title and professor of all lectures?

```
select Title, PersID as ProfID  
from Lecture;
```

select attributes or expressions of attributes



# Set Operations

SQL supports: **union, intersect, minus**

```
( select Name  
  from Researcher)
```

**union**

```
( select Name  
  from Professor);
```

```
( select Name  
  from Researcher)
```

**union all**

```
( select Name  
  from Professor);
```

# Grouping, Aggregation

Aggregate functions: **avg**, **max**, **min**, **count**, **sum**

```
select avg (Semester)  
from Student;
```

# Grouping, Aggregation

Aggregate functions: **avg, max, min, count, sum**

```
select avg (Semester)
```

```
from Student;
```

```
select PersID, sum (CP)
```

```
from Lecture
```

```
group by PersID;
```

# Grouping, Aggregation

Aggregate functions: **avg**, **max**, **min**, **count**, **sum**

```
select avg (Semester)
```

```
from Student;
```

having clause following group by

```
select PersID, sum (CP)
```

```
from Lecture
```

```
group by PersID;
```

```
select p.PersID, Name, sum (CP)  
  from Lecture l, Professor p  
 where l.PersID= p.PersID and level = 'FP'  
 group by p.PersID, Name  
 having avg (CP) >= 3;
```

# Imperative Processing in SQL

↙ Step 1:

**from** Lecture l, Professor p

**where** l.PersID= p.PersID **and** level = 'FP'

↙ Step 2:

**group by** p.PersID, Name

↙ Step 3:

**having avg** (CP)  $\geq$  3;

↙ Step 4:

**select** p.PersID, Name, **sum** (CP)



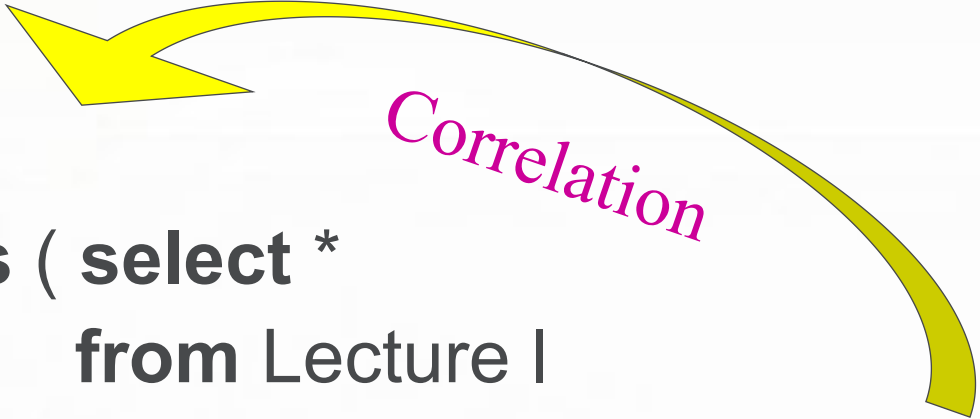
# Existential Quantification: exists sub-queries

```
select p.Name  
from Professor p  
where not exists ( select *  
                  from Lecture l  
                  where l.PersID = p.PersID );
```

EXIST R means R is not NULL

# Correlated Sub-queries

```
select p.Name  
from Professor p  
where not exists ( select *  
                  from Lecture l  
                  where l.PersID = p.PersID );
```



*Correlation*

cannot run without outer query

# Uncorrelated Sub-query

```
select Name  
from Professor           can run without outer query  
where PersID not in ( select PersID  
                           from Lecture);
```

Differences: Correlated v.s uncorrelated, comparing the above examples, which one is better?

More: <https://my.vertica.com/docs/4.1/HTML/Master/10374.htm>

# Sub-queries with all

```
select Name  
from Student  
where Semester >= all ( select Semester  
                        from Student);
```

Not as powerful as relational division!

# Subqueries in SELECT, FROM

```
select PersID, Name, ( select sum (CP) as load  from Lecture l
                        where p.PersID=l.PersID )
from Professor p;
```

```
select p.PersID, Name, l.load
from Professor p, ( select PersID, sum (CP) as load  from Lecture
                    group by PersID) l
where p.PersID = l.PersID;
```

Is this better than the simple Group By Query from before?



# Query Rewrite

```
select *  
from Researcher a  
where exists  
  (select *  
   from Professor p  
   where a.Supervisor= p.PersID and p.age < a.age);
```

⚡ Equivalent Join Query: Why is this better?

```
select a.*  
from Researcher a, Professor p  
where a.supervisor=p.PersID and p.age < a.age;
```

# Universal Quantification

- SQL does not support relational division directly
- Need to play tricks

$\{s \mid s \in \text{Student} \wedge$   
 $\forall l \in \text{Lecture} (l.\text{CP}=4 \Rightarrow$   
 $\exists a \in \text{attends} (a.\text{LecID}=l.\text{LecID} \wedge a.\text{StuID}=s.\text{StuID}))\}$

- Approach: Eliminate of  $\forall$  and  $\Rightarrow$

$$\forall t \in R (P(t)) = \neg(\exists t \in R(\neg P(t)))$$

$$R \Rightarrow T = \neg R \vee T$$



↙ Applying these rules:

$$\{s \mid s \in \text{Student} \wedge \neg (\exists l \in \text{Lecture} \neg (\neg (l.CP=4) \vee \exists a \in \text{attends}(a.LecID=l.LecID \wedge a.StuID=s.StuID)))\}$$

↙ Applying DeMorgan rules:

$$\{s \mid s \in \text{Student} \wedge \neg (\exists l \in \text{Lecture} (l.CP=4 \wedge \neg (\exists a \in \text{attends}(a.LecID=l.LecID \wedge a.StuID=s.StuID))))\}$$

↙ This can be implemented in SQL:

```
select *  
from Student s  
where not exists  
(select *  
from Lecture l  
where l.CP = 4 and not exists  
(select *  
from attends a  
where a.LecID = l.LecID and a.StuID=s.StuID) );
```

↙ This can be implemented in SQL:

**select \***

**from Student s**

**where not exists**

**(select \***

**from Lecture l**

**where l.CP = 4 and not exists**

**(select \***

**from attends a**

**where a.LecID = l.LecID and a.StuID=s.StuID) );**

Find students for which there is no  
four-hour lecture that  
student does not attend!

# Or do it this way

```
select a.StuID  
from attends a  
group by a.StuID  
having count (*) = (select count (*) from Lecture);
```



## Considering only 4 CP lectures

```
select a.StuID
from attends a, Lecture l
where a.LecID = l.LecID and l.CP = 4
group by a.StuID
having count (*) = (select count (*) from Lecture
                    where CP = 4);
```

# Null Values (NULL = UNKNOWN)

```
select count (*)
```

```
from Student
```

```
where Semester < 13 or Semester > =13;
```

**vs.**

```
select count (*)
```

```
from Student;
```

Are those two queries equivalent?

# Working with Null Values

1. **Arithmetics: Propagate **null**:** If an operand is null, the result is **null**.

`null + 1 -> null`

`null * 0 -> null`

2. **Comparisons: New Boolean value **unknown**.** All comparisons that involve a **null** value, evaluate to **unknown**.

`null = null -> unknown`

`null < 13 -> unknown`

`null > null -> unknown`

3. **Logic: Boolean operators are evaluated using the following tables (next slide):**

<b>not</b>	
<i>true</i>	false
<i>unknown</i>	unknown
<i>false</i>	true

<b>and</b>	<i>true</i>	<i>unknown</i>	<i>false</i>
<i>true</i>	true	unknown	false
<i>unknown</i>	unknown	unknown	false
<i>false</i>	false	false	false

<b>or</b>	<i>true</i>	<i>unknown</i>	<i>false</i>
<i>true</i>	true	true	true
<i>unknown</i>	true	unknown	unknown
<i>false</i>	true	unknown	false

4. **where:** Only tuples which evaluate to **true** are part of the query result. (**unknown** and **false** are equivalent here):

```
select count (*)  
from Student  
where Semester < 13 or Semester > =13;
```

5. **group by:** If exists, then there is a group for **null**.

```
select count (*)  
  
from Student  
group by Semester;
```

Predicates with null:

```
select count (*) from Student  
where Semester is null;
```

# Syntactic Sugar

**select \***

**from Student**

**where Semester > = 1 and Semester < = 6;**

**select \***

**from Student**

**where Semester between 1 and 6;**

**select \***

**from Student**

**where Semester in (2,4,6);**



# case

```
select StuID, ( case when Grade >= 5.5 then 'A'  
                when Grade >= 5.0 then 'B'  
                when Grade >= 4.5 then 'C'  
                when Grade >= 4.0 then 'D'  
                else 'F'end)
```

from tests;

- ⚡ Behaves like a switch: evaluate from top to bottom
- ⚡ No "break" needed because at most one clause executed.

# Comparisons with like

- ↙ "%" represents any sequence of characters (0 to n)
- ↙ "\_" represents exactly one character
- ↙ N.B.: For comparisons with = , % and \_ are normal chars.

**select \***

**from** Student

**where** Name **like** 'T%eophrastos';

**select distinct** Name

**from** Lecture l, attends a, Student s

**where** s.StuID= a.StuID **and** a.LecID = l.LecID  
**and** l.Title like '%thik%';

# Joins in SQL-92

- ↙ **cross join**: Cartesian product
- ↙ **natural join**:
- ↙ **join** or **inner join**: Theta-Join
- ↙ **left**, **right** or **full outer join**: outer join variants
- ↙ (union join: not discussed here)

```
select *  
  from  $R_1, R_2$   
  where  $R_1.A = R_2.B$ ;  
select *  
from  $R_1$  join  $R_2$  on  $R_1.A = R_2.B$ ;
```

# Left Outer Joins

```
select p.PersID, p.Name, t.PersID, t.Grade, t.StuID, s.StuID, s.Name
```

```
from Professor p left outer join
```

```
(tests t left outer join Student s  
on t.StuID= s.StuID)
```

```
on p.PersID=t.PersID;
```

PersID	p.Name	t.PersID	t.Grade	t.StuID	s.StuID	s.Name
2126	Russel	2126	1	28106	28106	Carnap
2125	Sokrates	2125	2	25403	25403	Jonas
2137	Kant	2137	2	27550	27550	Schopen- hauer
2136	Curie	-	-	-	-	-
:	:	:	:	:	:	:

# Right Outer Joins

```
select p.PersID, p.Name, t.PersID, t.Grade, t.StuID, s.StuID, s.Name
from Professor p right outer join
    (tests t right outer join Student s on
        t.StuID= s.StuID)
on p.PersID=t.PersID;
```

PersID	p.Name	t.PersID	t.Grade	t.StuID	s.StuID	s.Name
2126	Russel	2126	1	28106	28106	Carnap
2125	Sokrates	2125	2	25403	25403	Jonas
2137	Kant	2137	2	27550	27550	Schopen- hauer
-	-	-	-	-	26120	Fichte
⋮	⋮	⋮	⋮	⋮	⋮	⋮

# Full Outer Joins

```
select p.PersID, p.Name, t.PersID, t.Grade, t.StuID,  
s.StuID, s.Name
```

```
from Professor p full outer join
```

```
(tests t full outer join Student s on
```

```
t.StuID= s.StuID)
```

```
on p.PersID=t.PersID;
```



p.PersID	p.Name	t.PersID	t.Grade	t.StuID	s.StuID	s.Name
2126	Russel	2126	1	28106	28106	Carnap
2125	Sokrates	2125	2	25403	25403	Jonas
2137	Kant	2137	2	27550	27550	Schopen- hauer
-	-	-	-	-	26120	Fichte
⋮	⋮	⋮	⋮	⋮	⋮	⋮
2136	Curie	-	-	-	-	-
⋮	⋮	⋮	⋮	⋮	⋮	⋮

# **connect by Clause (Oracle)**

```
select Title
from Lecture
where LecID in (select prerequisite
                from requires
                connect by follow-up = prior prerequisite
                start with follow-up = (select LecID
                                        from Lecture
                                        where Title = ... ));
```

# Recursion in DB2/SQL99

**with** TransLecture (First, Next)

**as** (**select** prerequisite, follow-up **from** requires  
**union all**

**select** t.First, r.follow-up

**from** TransLecture t, requires r

**where** t.Next= r.prerequisite)

**select** Title **from** Lecture **where** LecID in

(**select** First **from** TransLecture **where** Next in

(**select** LecID **from** Lecture

**where** Title = 'DMD') )

# Data Manipulation Language

Insert tuples

**insert into** attends

**select** StuID, LecID

**from** Student, Lecture

**where** Title= `DMD`;

**insert into** Student (StuID, Name)

**values** (28121, `Archimedes`);

# Deletion of tuples, Update

**delete** Student

**where** Semester > 13;

**update** Student

**set** Semester = Semester + 1;

# Snapshot Semantics

1. Phase 1: mark tuples which are affected by the update
2. Phase 2: implement update on marked tuples

Otherwise, indeterministic execution of updates:

**delete from** requires

**where** prerequisite **in** (**select** follow-up  
**from** requires);

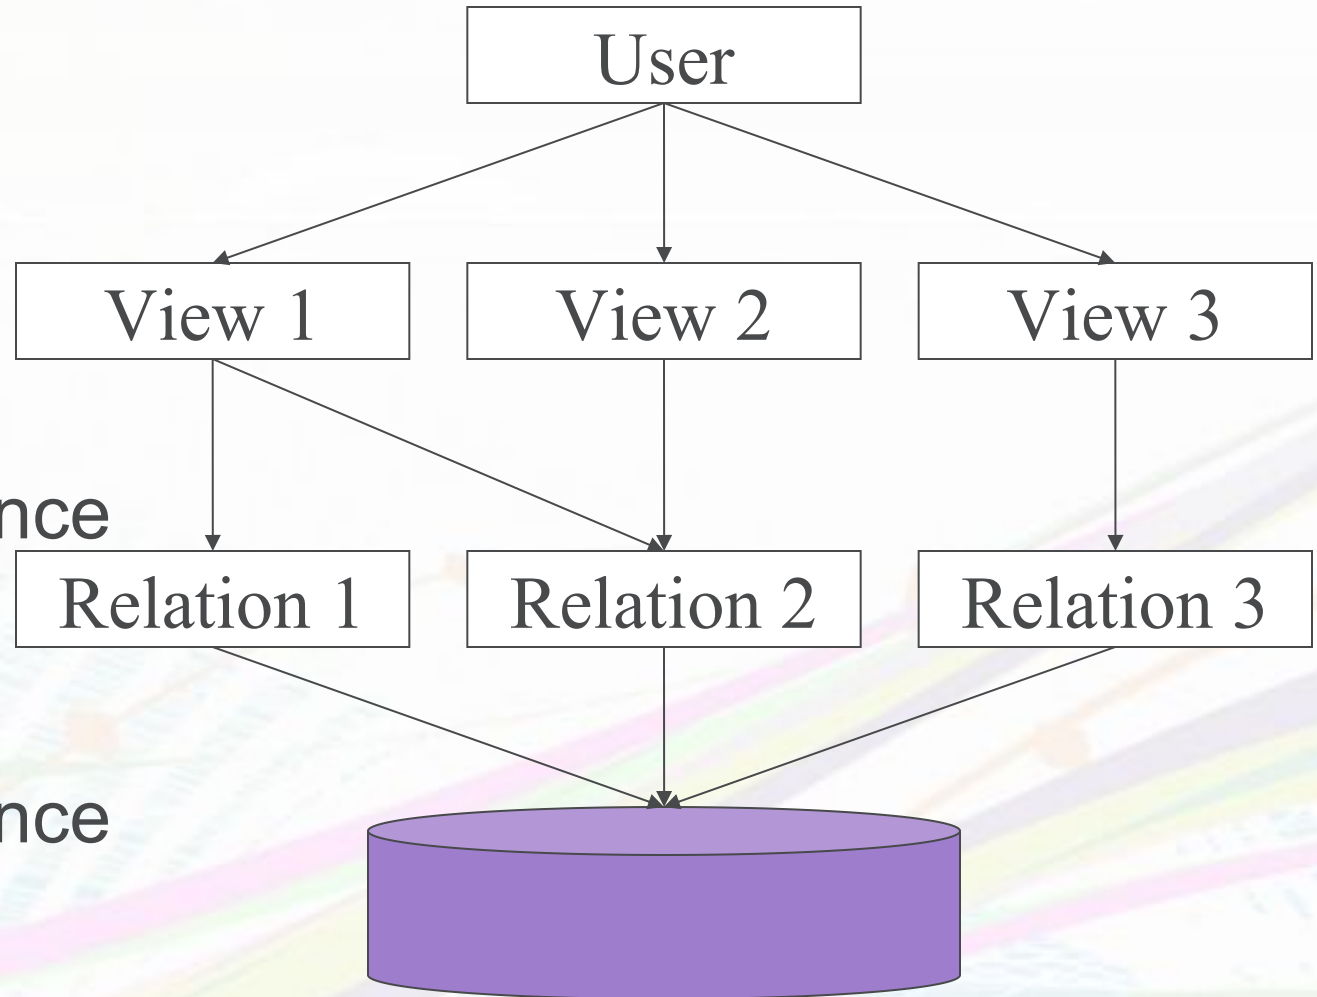


requires	
Prerequisite	Follow-up
5001	5041
5001	5043
5001	5049
5041	5216
5043	5052
5041	5052
5052	5229

**delete from requires**

**where Prerequisite in (select Follow-up  
from requires);**

# Views for Logical Data Independence



# Views ...

for privacy

```
create view testView as  
  select StuID, LecID, PersID  
  from tests;
```

# Views ...

**for simpler queries**

**create view StudProf (Sname, Semester, Title, Pname)**  
**as**

```
select s.Name, s.Semester, l.Title, p.Name  
from Student s, attends a, Lecture l, Professor p  
where s.StuID=a.StuID and a.LecID=l.LecID and  
l.PersID= p.PersID;
```

```
select distinct Semester  
from StudProf  
where PName='Someone';
```

# Views for is-a relationships

```
create table Employee
```

```
  (PersID integer not null,
```

```
   Name   varchar (30) not null);
```

```
create table ProfData
```

```
  (PersID integer not null,
```

```
   Level   character(2),
```

```
   Room    integer);
```

```
create table ResearcherData
```

```
  (PersID   integer not null,
```

```
   area      varchar(30),
```

```
   Supervisor integer);
```

**create view Professor as**

**select \***

**from Employee e, ProfData d**

**where e.PersID=d.PersID;**

**create view Researcher as**

**select \***

**from Employee e, ResearcherData d**

**where e.PersID=d.PersID;**

**➔ Subtypes implemented as a view**

**create table Professor**

(PersID      **integer not null,**  
Name        **varchar (30) not null,**  
Level       **character (2),**  
Room        **integer);**

**create table Researcher**

(PersID      **integer not null,**  
Name        **varchar (30) not null,**  
area         **varchar (30),**  
Supervisor **integer);**

**create table OtherEmps**

(PersID      **integer not null,**  
Name        **varchar (30) not null);**



```
create view Employee as  
  (select PersID, Name  
   from Professor)  
  union  
  (select PersID, Name  
   from Researcher)  
  union  
  (select*  
   from OtherEmps);
```

➔ Supertypes implemented as a view

# Updatable Views

## Example view which is not updatable

**create view ToughProf (PersID, AvgGrade) as**

**select PersID, avg(Grade)**

**from tests**

**group by PersID;**

**update ToughProf set AvgGrade= 6.0**

**where PersID = 4711;**

**insert into ToughProf**

**values (4711, 6.0);**

**SQL tries to avoid indeterminisms.**

# What about this?

```
create view ToughProf (PersID, AvgGrade) as  
select PersID, avg(Grade)  
from tests  
group by PersID;
```

```
delete ToughProf  
where PersID = 4711;
```

# Views and Updates

## Example view which is not updatable

create view LectureView as

```
select Title, CP, Name
```

```
from Lecture l, Professor p
```

```
where l.PersID = p.PersID;
```

insert into LectureView

```
values ('Nihilismus', 2, 'Nobody');
```

There are scenarios in which the "insert" is meaningful.  
There are scenarios in which SQL would have to guess.  
SQL is conservative and does not allow any scenario.

# Views and Updates in SQL

## ↙ A SQL view is updatable iff

The view involves only one base relation

The view involves the key of that base relation

The view does NOT involve aggregates, group by, or duplicate-elimination

