

# Constraints & Triggers

The background of the slide is a complex, abstract composition. It features a light gray grid of thin, curved lines that sweep across the frame. Overlaid on this are several thick, vibrant lines in shades of magenta, lime green, and dark purple. These lines are dynamic, with some following straight paths while others curve. Small, solid-colored squares (orange, green, and purple) are placed at various points along these lines, acting as markers or data points. The overall effect is one of movement and interconnectedness, suggesting a network or a system of relationships.

# Why use a DBMS? (Week 1)

- Avoid redundancy and **inconsistency**
- Rich (declarative) access to the data
- Synchronize concurrent data access
- Recovery after system failures
- Security and privacy
- Reduce cost and pain to do something useful
  - There is always an alternative!!!

# Integrity of Data

## ↙ Example Constraints

Keys

attribute domains

Referential integrity (foreign keys  $\rightarrow$  keys)

...

## ↙ Static Constraints

Constraints that any instance of a DB must meet

## ↙ Dynamic Constraints

Constraints on a state transition of the DB

# Who checks? DB vs. App

## ⚡ Why implement constraints in the DB?

Good way to annotate & document schema

DB is a central point (once and for all cases)

Safety: in case you forget it in the app

Useful for DB-level optimization

- Constraint: all students are older than 18 years.
- Query: `SELECT * FROM Student WHERE age < 17;`
- Query can be evaluated without looking at any student.

## ⚡ Why implement constraints in the App?

Meaningful error messages.

## ⚡ It is important to do both!!!



# Referential Integrity Constraints

## Foreign Keys

- ↙ Refer to tuple from a different relation
- ↙ E.g., PersID in table Lecture refers to Professor

## Definition: Referential Integrity

For every foreign key one of the two conditions must hold

the value of the foreign key is *NULL* or  
the referenced tuple must exist

# Referential Integrity in SQL

## SQL Syntax to declare keys and foreign keys:

Key: **unique**

Primary key: **primary key**

Foreign key: **foreign key**

## Example:

**create table  $R$**

**(  $\alpha$  integer primary key,**

**$\beta$  varchar(30) unique,**

**... );**

**create table  $S$**

**( ...,**

**$\kappa$  integer references  $R$  );**

# Maintaining referential integrity?

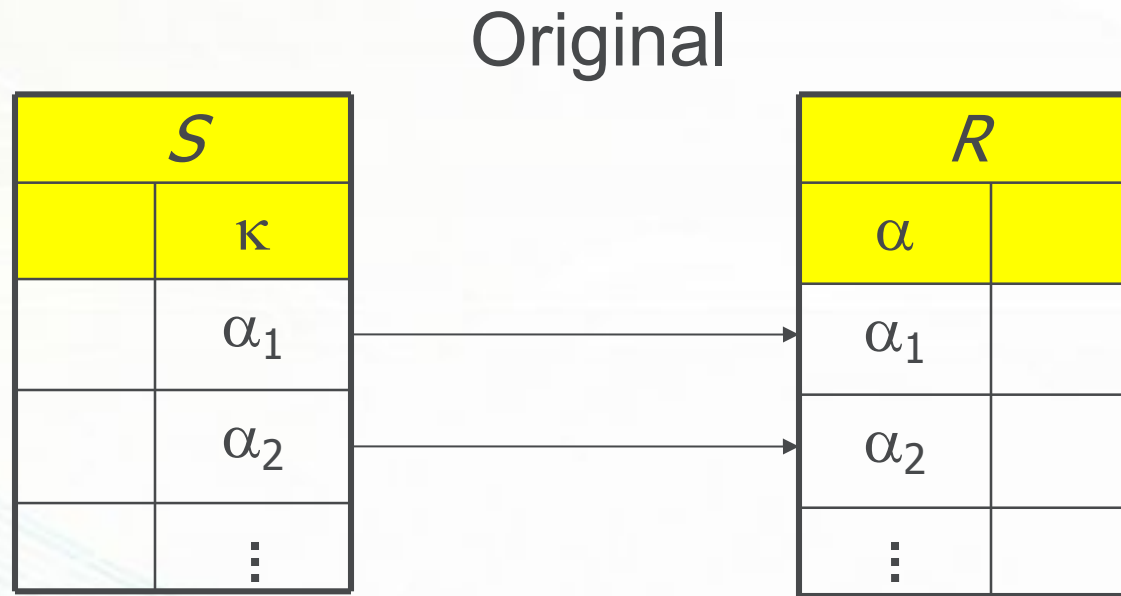
Updates of referenced data which result in a violation

1. Default: reject the update (return an error)
2. **cascade**: propagate update
3. **set null**: set references to null
4. (Set references to default value. Not supported in SQL.)

The right choice depends on the ER model and operations.

- ✚ e.g. weak vs. strong entities
- ✚ relations that implement N:M relationships
- ✚ 1:N relations
- ✚ Exercise: extend rules for ER->relational translation!

# Maintaining referential integrity



Update

update  $R$

**set**  $\alpha = \alpha'_1$


**where**  $\alpha = \alpha_1$ ;

delete from  $R$


**where**  $\alpha = \alpha_1$ ;



# Cascade (weak entities, n:m relationships)



Update of  $S$



Delete in  $S$

**create table  $S$**

( ...,

$\kappa$  integer references  $R$

on update cascade );


**create table  $S$**

( ...,


$\kappa$  integer references  $R$

on delete cascade );

## Set Null (strong entities)



Update of  $S$



Update of  $S$

**create table  $S$**

**( ...,**

**$\kappa$  integer references  $R$**

**on update set null );**

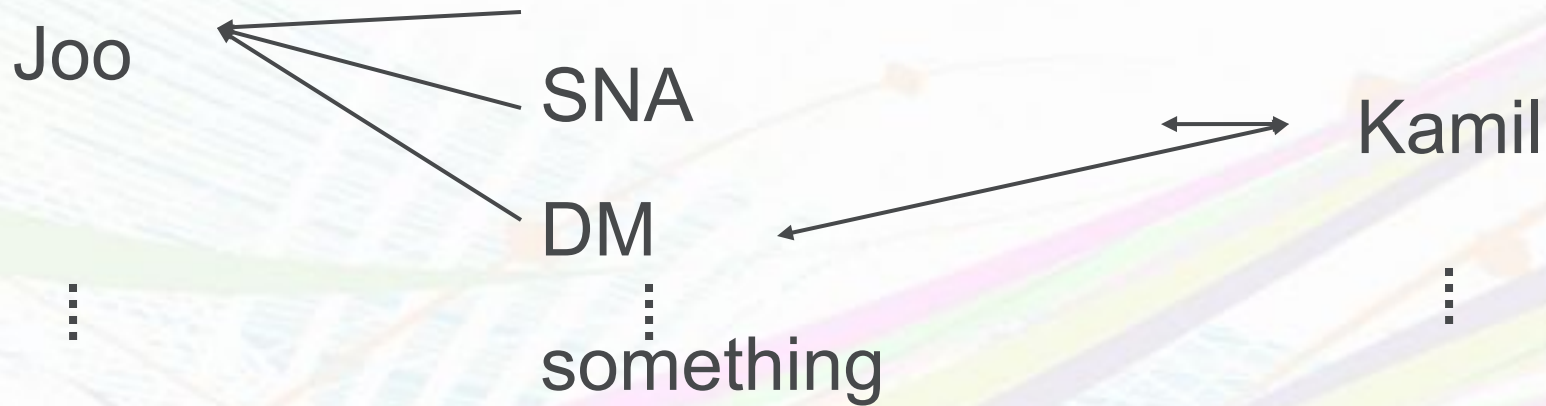

**create table  $S$**

**( ...,**

**$\kappa$  integer references  $R$**

**on delete set null );**

# Cascading Deletes



**create table Lecture**

**( ...,**

**PersID integer**

**references Professor**

**on delete cascade);**

**create table attends**

**( ...,**

**ID integer**

**references Lecture**

**on delete cascade);**

# Constraints on Attribute Domains

## ↙ Integer domains

... **check** Semester **between** 1 and 13

## ↙ Enum types

... **check** Level **in** ('Assistant', 'Associate', 'Full') ...



# Uni-DB schema with Constraints

**create table Student**

( StuID      **integer primary key,**  
Name        **varchar(30) not null,**  
Semester   **integer check (Semester between 1 and 13));**

**create table Professor**

( PersID     **integer primary key,**  
Name        **varchar(30) not null,**  
Level        **character(2) check (Level in ('AP','CP','FP')),**  
Room        **integer unique );**

## **create table Assistant**

( PersID	<b>integer primary key,</b>
Name	<b>varchar(30) not null,</b>
Area	<b>varchar(30),</b>
Boss	<b>integer,</b>
<b>foreign key</b>	<b>(Boss) references Professor</b> <b>on delete set null);</b>

## **create table Lecture**

( LecID	<b>integer primary key,</b>
Title	<b>varchar(30),</b>
CP	<b>integer,</b>
PersID	<b>integer references Professor</b> <b>on delete set null);</b>

**create table** attends

```
( StuID          integer references Student
                        on delete cascade,
  LecID          integer references Lecture
                        on delete cascade,
  primary key    (StuID, LecID));
```

**create table** requires

```
( Prerequisite integer references Lecture
                        on delete cascade,
  Follow-up    integer references Lecture
                        on delete cascade,
  primary key   (Prerequisite, Follow-up));
```

## **create table tests**

```
( StuID      integer references Student
              on delete cascade,
  LecID       integer references Lecture,
  PersID      integer references Professor
              on delete set null,
  Grade       numeric (3,2)
              check (Grade between 1.0 and 6.0),
  primary key (StuID, LecID));
```

# constraints on tuple

**create table** Student

( StuID        **integer primary key,**

Name        **varchar(30) not null,**

Semester    **integer,**

**check** (Semester **between** 1 and 13)

);



# 1:1 Relationships (Wedding)

```
create table Man(  
    name  varchar(30) primary key;  
    spouse varchar(30) references Woman);  
  
create table Woman(  
    name  varchar(30) primary key;  
    spouse varchar(30) references Man);
```

⚡ Legal: Helga marries Hugo, but Hugo does not marry Helga.

Mutual marriage cannot be expressed in SQL.

How would you model marriage in SQL?

(cha 7.1.3)

⚡ N.B.: The real implementation is based on **transactions!**

# Trigger (ECA Rules)

**create trigger noDegradation**  
**before update on Professor**  
**for each row (or statement)**  
**when (old.Level is not null)**  
**begin**

**if :old.Level = 'Associate' and :new.Level = 'Assistant' then**  
**:new.Level := 'Associate';**

**end if;**

**if :old.Level = 'Full' then**  
**:new.Level := 'Full'**

**end if;**

**if :new.Level is null then**  
**:new.Level := :old.Level;**  
**end if;**

**end**

**Event**

**Condition**

**Action**

# Dangers of Triggers

```
create trigger weddingMan
after update on Man
for each row
when (true)
begin
    update Woman set spouse = :new.Name
    where name = :new.spouse;
    update Woman set spouse = null
    where name = :old.spouse;
end
```

What happens if we write a weddingWoman trigger?  
Is marriage better modeled statically or dynamically?



# Exercise

- ↙ 1. What is theta join?
- ↙ 2. Write an example of theta join.

Template:

BS1#1, First\_name Last\_name

Solution\_1:.....