

Academic Honesty Policy

- ↙ Collaboration **is strongly encouraged**, but...
- ↙ Each student must write down his/her own solutions independently in his or her own words.
- ↙ Each student must submit a list of anyone with whom the assignment was discussed.
- ↙ All sources (internet included) must be cited and credited.
- ↙ Solution sets from this course or any other course may not be used under any circumstances.

↙ Mid-term Evaluation

- Participation
- Assignments
- Project Phase I

↙ Mid-term Exam

- Lecture time on Nov 16th, 2015

↙ Course Feedback and Discussion

- One representative from each seminar group
- Office 403, 11:00–12:00 Sep 25

Again, what if you have questions/problems?

- ↙ Slides, Textbooks, and Google
- ↙ Your classmates
- ↙ Piazza
- ↙ Your seminar instructors
- ↙ I am always available on Monday for your questions (403, 415, or around)



Again, how to succeed in the course?

- ↙ Read Slides, Textbooks, and use Google
- ↙ Work with your classmates
- ↙ Ask questions and help others on Piazza
- ↙ Discuss with your seminar instructors
- ↙ Participate lectures/seminars, handin assignments, handin projects



Functional Dependencies

The background of the slide is an abstract composition. It features a light gray grid of thin, curved lines that sweep across the frame. Overlaid on this are several thick, vibrant lines in shades of magenta, lime green, and dark purple. These lines are mostly diagonal, trending from the bottom-left towards the top-right. There are also some thinner, curved lines in orange and green, some of which have small square markers at their ends, giving the impression of a dynamic or data-driven environment.

Relational Design Theory

⚡ How to assess the quality of a schema

redundancy

integrity constraints

(Formal properties of a schema) Quality seal: normal forms

⚡ How to improve the quality of a schema

synthesis algorithm

decomposition algorithm

⚡ How to construct a (high-quality) schema

start with universal relation

apply synthesis or decomposition algorithms

Why not redundancy?

↙ Waste of storage space

importance is diminishing as storage gets cheaper
(disk density will even increase in the future)

↙ Additional work to keep multiple copies of data consistent

multiple updates in order to accomodate one event

↙ Additional code to keep multiple copies of data consistent

Somebody needs to implement the logic

but..

- ↙ Sometimes redundancy is necessary.
- ↙ possible to improve data locality
- ↙ Space (memory, disk) is no longer the problem it used to be => tradeoff space/performance
- ↙ Fault tolerance
- ↙ We will discuss this later. Now we will focus on avoiding redundancy in a schema

Functional Dependencies

↙ $X \rightarrow A$ is an assertion about a relation R that whenever two tuples of R agree on all the attributes of X , then they must also agree on the attribute A .

➤ Say “ $X \rightarrow A$ holds in R .”

➤ **Convention:** ..., X, Y, Z represent sets of attributes; A, B, C, \dots represent single attributes.

➤ **Convention:** no set formers in sets of attributes, just ABC , rather than $\{A, B, C\}$.

Example

Drinkers(name, addr, beersLiked, manf, favBeer)

↙ Reasonable FD's to assert:

1. name \rightarrow addr
2. name \rightarrow favBeer
3. beersLiked \rightarrow manf

Example Data

name	addr	beersLiked	manf	favBeer
Janeway	Voyager	Bud	A.B.	WickedAle
Janeway	Voyager	WickedAle	Pete's WickedAle	
Spock	Enterprise	Bud	A.B.	Bud

Because name -> addr

Because name -> favBeer

Because beersLiked -> manf

FD's With Multiple Attributes

- ✧ No need for FD's with > 1 attribute on right.

But sometimes convenient to combine FD's as a shorthand.

Example: $\text{name} \rightarrow \text{addr}$ and

$\text{name} \rightarrow \text{favBeer}$ become

$\text{name} \rightarrow \text{addr favBeer}$

- ✧ > 1 attribute on left may be essential.

Example: $\text{bar beer} \rightarrow \text{price}$

Keys of Relations

- ↙ K is a *superkey* for relation R if K functionally determines all of R .
- ↙ K is a *key* for R if K is a superkey, but no proper subset of K is a superkey.

Example

Drinkers(name, addr, beersLiked, manf, favBeer)

⚡ {name, beersLiked} is a superkey because together these attributes determine all the other attributes.

name \rightarrow addr favBeer

beersLiked \rightarrow manf

Example, Cont.

- ↙ $\{\text{name}, \text{beersLiked}\}$ is a **key** because neither $\{\text{name}\}$ nor $\{\text{beersLiked}\}$ is a superkey.

name doesn't \rightarrow manf ; beersLiked
doesn't \rightarrow addr .

- ↙ There are no other keys, but lots of superkeys.

Any superset of $\{\text{name}, \text{beersLiked}\}$.

E/R and Relational Keys

- ✚ Keys in E/R concern **entities**.
- ✚ Keys in relations concern **tuples**.
- ✚ Usually, one tuple corresponds to one entity, so the ideas are the same.
- ✚ But --- in poor relational designs, one entity can become several tuples, so E/R keys and Relational keys are different.

Example Data

name	addr	beersLiked	manf	favBeer
Janeway	Voyager	Bud	A.B.	WickedAle
Janeway	Voyager	WickedAle	Pete's WickedAle	
Spock	Enterprise	Bud	A.B.	Bud

Relational key = {name beersLiked}

But in E/R, name is a key for Drinkers, and beersLiked is a key for Beers.

Note: 2 tuples for Janeway entity and 2 tuples for Bud entity.

Where Do Keys Come From?

1. Just assert a key K .

The only FD's are $K \rightarrow A$ for all attributes A .

2. Assert FD's and deduce the keys by systematic exploration.

E/R model gives us FD's from entity-set keys and from many-one relationships.

More FD's From “Physics”

- ↙ Example: “no two courses can meet in the same room at the same time” tells us: hour room → course.

Inferring FD's

↙ We are given FD's $X_1 \rightarrow A_1, X_2 \rightarrow A_2, \dots, X_n \rightarrow A_n$, and we want to know whether an FD $Y \rightarrow B$ must hold in any relation that satisfies the given FD's.

Example: If $A \rightarrow B$ and $B \rightarrow C$ hold, $A \rightarrow C$ holds. (**Transitivity**)

↙ Important for design of good relation schemas.

Inference Test

↙ To test if $Y \rightarrow B$, start by assuming two tuples agree in all attributes of Y .

← Y →

0000000...0

00000?? ... ?

Inference Test – (2)


↙ Use the given FD's to infer that these tuples must also agree in certain other attributes.

If B is one of these attributes, then $Y \rightarrow B$ is true.

Otherwise, the two tuples, with any forced equalities, form a two-tuple relation that proves $Y \rightarrow B$ does not follow from the given FD's.

Closure Test

- ↙ An easier way to test is to compute the *closure* of Y , denoted Y^+ .
- ↙ **Basis**: $Y^+ = Y$.
- ↙ **Induction**: Look for an FD left side X that is a subset of the current Y^+ . If the FD is $X \rightarrow A$, add A to Y^+ .



Example

↙ A relation with attributes ABCDEF

↙ FD's: $A, B \rightarrow C$, $BC \rightarrow AD$, $D \rightarrow E$

↙ AB^+

Finding All Implied FD's


↙ **Motivation:** “normalize” the process where we break a relation schema into two or more schemas.

↙ Example: $ABCD$ with FD's $AB \rightarrow C$, $C \rightarrow D$, and $D \rightarrow A$.

Decompose into ABC , AD . What FD's hold in ABC ?

Not only $AB \rightarrow C$, but also $C \rightarrow A$!

Why?



Thus, tuples in the projection with equal C's have equal A's;
 $C \rightarrow A$.

Basic Idea

1. Start with given FD's and find all *nontrivial* FD's that follow from the given FD's.

Nontrivial = left and right sides disjoint.

2. Restrict to those FD's that involve only attributes of the projected schema.

Simple, Exponential Algorithm

1. For each set of attributes X , compute X^+ .
2. Add $X \rightarrow A$ for all A in $X^+ - X$.
3. However, drop $XY \rightarrow A$ whenever we discover $X \rightarrow A$.
 - ◆ Because $XY \rightarrow A$ follows from $X \rightarrow A$ in any projection.
4. Finally, use only FD's involving projected attributes.

A Few Tricks

- ↙ No need to compute the closure of the empty set or of the set of all attributes.
- ↙ If we find $X^+ = \text{all attributes}$, so is the closure of any superset of X .

Example

↙ ABC with FD's $A \rightarrow B$ and $B \rightarrow C$. Project onto AC .

$A^+ = ABC$; yields $A \rightarrow B$, $A \rightarrow C$.

- We do not need to compute AB^+ or AC^+ .

$B^+ = BC$; yields $B \rightarrow C$.

$C^+ = C$; yields nothing.

$BC^+ = BC$; yields nothing.

Example --- Continued

↙ Resulting FD's: $A \rightarrow B$, $A \rightarrow C$, and $B \rightarrow C$.


↙ Projection onto AC : $A \rightarrow C$.

Only FD that involves a subset of $\{A, C\}$.

A Geometric View of FD's

- ↙ Imagine the set of all *instances* of a particular relation.
- ↙ That is, all finite sets of tuples that have the proper number of components.
- ↙ Each instance is a point in this space.

Example: $R(A,B)$




An FD is a Subset of Instances

- ↙ For each FD $X \rightarrow A$ there is a subset of all instances that satisfy the FD.
- ↙ We can represent an FD by a region in the space.
- ↙ Trivial FD = an FD that is represented by the entire space.

Example: $A \rightarrow A$.

Example: $A \rightarrow B$ for $R(A,B)$




Representing Sets of FD's

- ↙ If each FD is a set of relation instances, then a collection of FD's corresponds to the intersection of those sets.

Intersection = all instances that satisfy all of the FD's.

Example

Instances satisfying
 $A \rightarrow B$, $B \rightarrow C$, and
 $CD \rightarrow A$



Implication of FD's

⚡ If an FD $Y \rightarrow B$ follows from FD's $X_1 \rightarrow A_1, \dots, X_n \rightarrow A_n$, then the region in the space of instances for $Y \rightarrow B$ must include the intersection of the regions for the FD's $X_i \rightarrow A_i$.

That is, every instance satisfying all the FD's $X_i \rightarrow A_i$ surely satisfies $Y \rightarrow B$.

But an instance could satisfy $Y \rightarrow B$, yet not be in this intersection.

Example

