


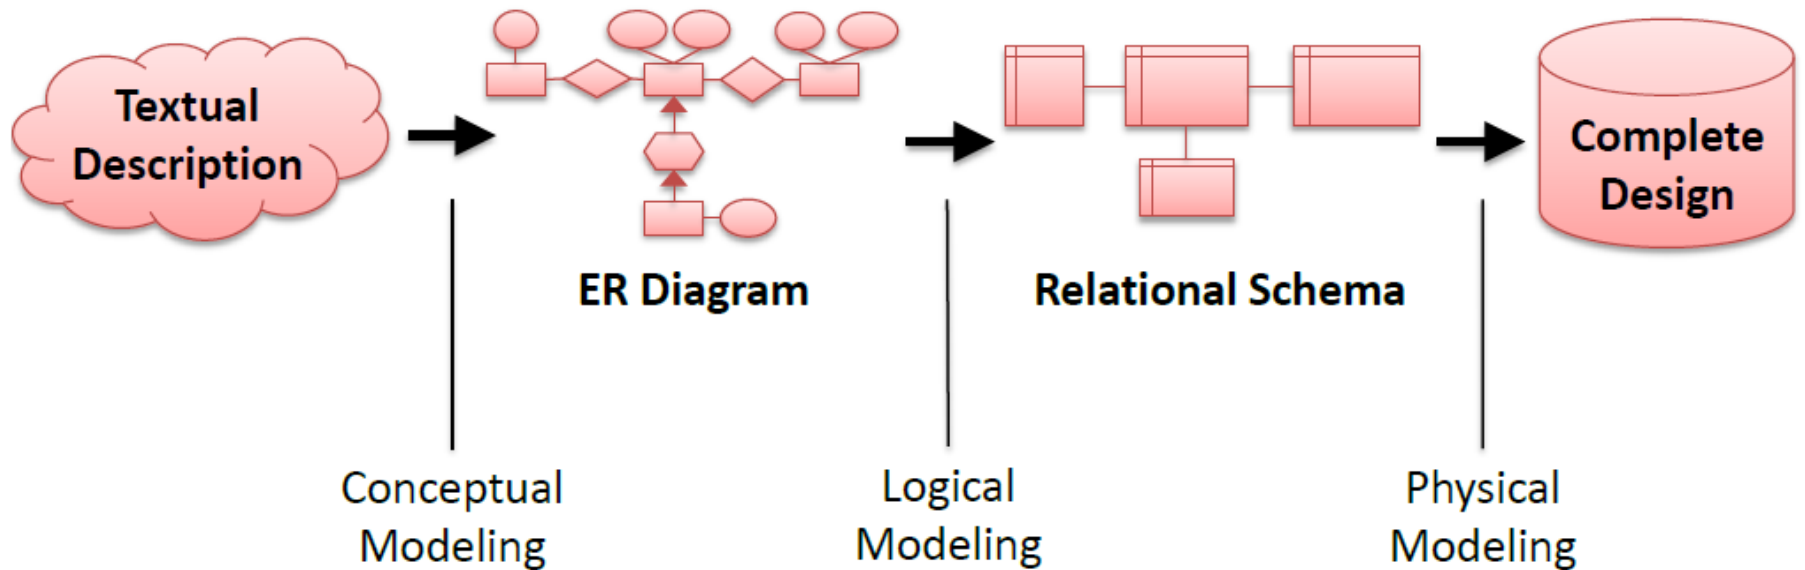
File management, Hash & External sorting

Monday, Oct. 12, 2015






Database Design



Under
the
hood
of
DBMS




Leverage OS for disk/file management?

- Layers of abstraction are good ... but:
 - Unfortunately, OS often **gets in the way** of DBMS
- DBMS wants/needs to do things “its own way”
 - Specialized prefetching
 - Control over buffer replacement policy
 - Control over thread/process scheduling
 - Arises when OS scheduling conflicts with DBMS locking
 - Control over flushing data to disk

Disks and Files

- DBMS stores information on disks. (Really?)
 - but: disks are (relatively) VERY slow!
- Major implications for DBMS design:
 - **READ:** disk -> main memory (RAM).
 - **WRITE:** reverse
 - Both are high-cost operations, relative to in-memory operations, so must be planned carefully!

The Storage Hierarchy




Anatomy of a Disk

Unlike RAM, time to retrieve a disk page varies depending upon location on disk.

relative placement of pages on disk is important!

- Sector
- Track
- Cylinder
- Platter
- Block size = multiple of sector size (which is fixed)




Accessing a Disk Page

- Time to access (read/write) a disk block:
 - *seek time*: moving arms to position disk head on track
 - *rotational delay*: waiting for block to rotate under head
 - *transfer time*: actually moving data to/from disk surface




Rotational Delay



Accessing a Disk Page

- Relative times?
 - *seek time*: about 1 to 20msec
 - *rotational delay*: 0 to 10msec
 - *transfer time*: < 1msec per 4KB page




Rules of thumb...

Memory access much faster than disk I/O (~ 1000x)

“Sequential” I/O faster than “random” I/O (~ 10x)

Disk Arrays: RAID

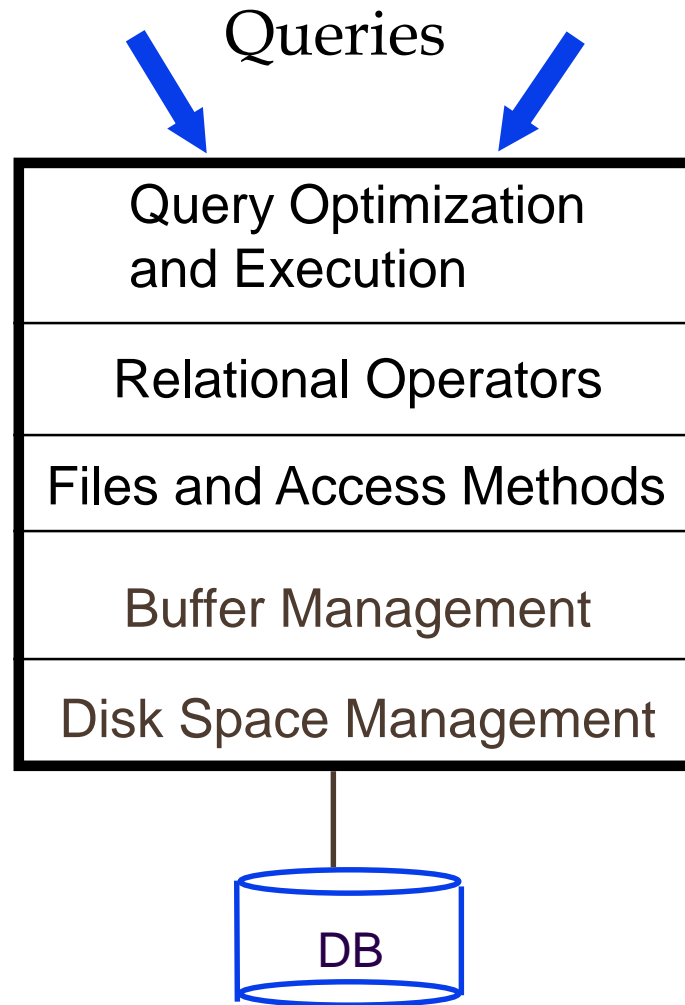


- Benefits:
 - Higher throughput (via data “striping”)
 - Longer MTTF (via redundancy)

Disk Space Management


- Lowest layer of DBMS software manages space on disk
- Higher levels call upon this layer to:
 - allocate/de-allocate a page
 - read/write a page
- Best if requested pages are stored sequentially on disk! Higher levels don't need to know if/how this is done, nor how free space is managed.

DBMS Layers

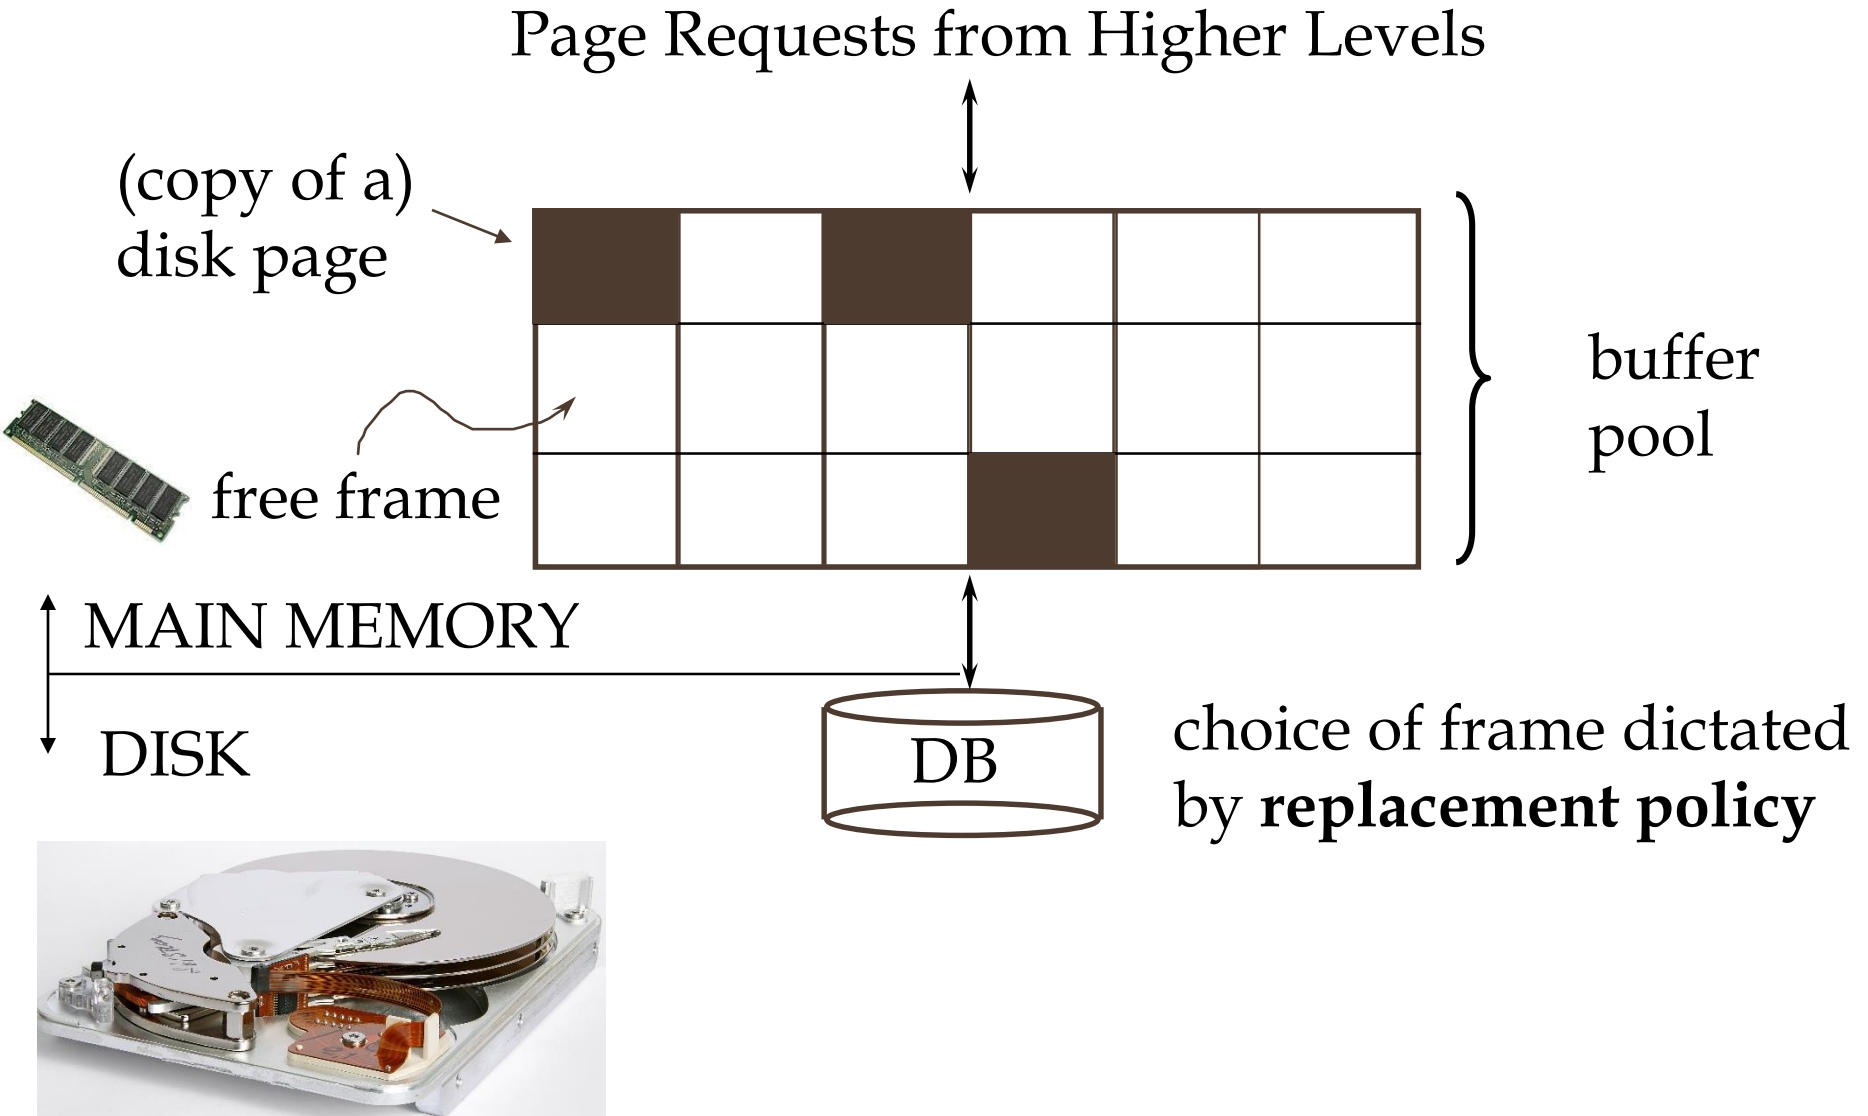


Buffer Management in a DBMS

- Data must be in RAM for DBMS to operate on it!
- Buffer Mgr hides the fact that not all data is in RAM



Buffer Management in a DBMS



Buffer Replacement Policy

- Frame is chosen for replacement by a *replacement policy*:
 - Least-recently-used (LRU), MRU, Clock, etc.
- Policy -> big impact on # of I/O 's; depends on the *access pattern*.



LRU Replacement Policy

- *Least Recently Used (LRU)*
 - replace the frame which has the oldest (earliest) time
 - very common policy: intuitive and simple
- Problems?
- *# buffer frames < # pages in file* means each page request causes an I/O. MRU much better in this situation (but not in all situations, of course).




Sequential Flooding – Illustration

LRU:

BUFFER POOL			
102	116	242	105

MRU:

BUFFER POOL			
102	116	242	105



Repeated scan of file ...


Sequential Flooding – Illustration

LRU:

BUFFER POOL			
1	116	242	105

MRU:

BUFFER POOL			
102	116	242	105



Repeated scan of file ...


Sequential Flooding – Illustration

LRU:

BUFFER POOL			
1	2	242	105

MRU:

BUFFER POOL			
102	116	242	105



Repeated scan of file ...


Sequential Flooding – Illustration

LRU:

BUFFER POOL			
1	2	3	105

MRU:

BUFFER POOL			
102	116	242	105



Repeated scan of file ...


Sequential Flooding – Illustration

LRU:

BUFFER POOL			
1	2	3	4

MRU:

BUFFER POOL			
102	116	242	105



Repeated scan of file ...

Sequential Flooding – Illustration


LRU:

BUFFER POOL			
1	2	3	4

will not re-use these pages;

MRU:

BUFFER POOL			
102	116	242	105



Repeated scan of file ...

Sequential Flooding – Illustration


LRU:

BUFFER POOL			
1	2	3	4

will not re-use these pages;

MRU:

BUFFER POOL			
1	116	242	105



Repeated scan of file ...

Sequential Flooding – Illustration


LRU:

BUFFER POOL			
1	2	3	4

will not re-use these pages;

MRU:

BUFFER POOL			
2	116	242	105



Repeated scan of file ...

Sequential Flooding – Illustration


LRU:

BUFFER POOL			
1	2	3	4

will not re-use these pages;

MRU:

BUFFER POOL			
3	116	242	105



Repeated scan of file ...

Sequential Flooding – Illustration


LRU:

BUFFER POOL			
1	2	3	4

will not re-use these pages;

MRU:

BUFFER POOL			
4	116	242	105



Repeated scan of file ...


Other policies?

- LRU is often good - but needs timestamps and sorting on them
- something easier to maintain?


“Clock” Replacement Policy

Main ideas:


- Approximation of LRU.
- Instead of maintaining & sorting time-stamps, find a ‘reasonably old’ frame to evict.
- How? by round-robin, and marking each frame - frames are evicted the second time they are visited.




“Clock” Replacement Policy



“Clock” Replacement Policy




“Clock” Replacement Policy



Files


- FILE: A collection of pages, each containing a collection of records.
- Must support:
 - insert/delete/modify record
 - read a particular record (specified using *record id*)
 - scan all records (possibly with some conditions on the records to be retrieved)

Heap File Using Lists



- Any problems?

Heap File Using a Page Directory




Page Formats

- fixed length records
- variable length records

Problem definition

Q: How would you store records on a page/file, such that

1. you can point to them
2. you can insert/delete records with few disk accesses



Page Formats

Important concept: *rid* == record id

Q0: why do we need it?


A0: eg., for indexing

Q1: How to mark the location of a record?

A1: rid = record id = page-id & slot-id


Q2: Why not its byte offset in the file?

A2: too much re-organization on ins/del.




Fixed length records

- Q: How would you store them on a page/file?




Fixed length records

- OK – how about insertion?




Fixed length records

- OK – how about insertion?




Fixed length records

- How about deletion?




Fixed length records

- Q: How would you store them on a page/file?
- A2: Bitmaps



Fixed length records

- Q: How would you store them on a page/file?
- A2: Bitmaps : ✓ insertions, ✓ deletions



Variable length records

- Q: How would you store them on a page/file?

occupied records

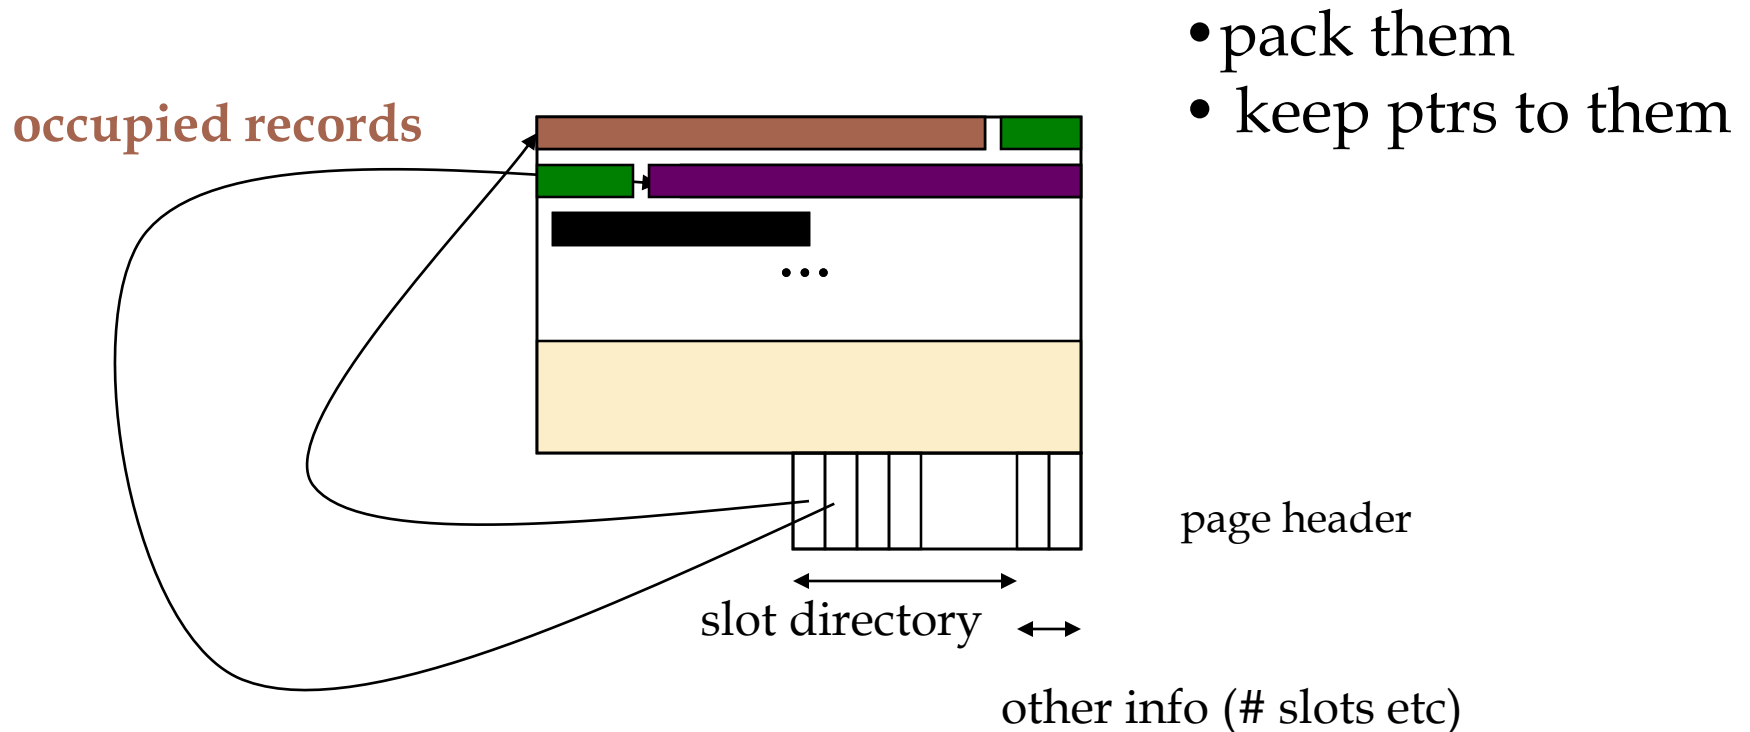


page header

SLOTTED PAGE

Variable length records


- Q: How would you store them on a page/file?



Variable length records

- Q: How would you store them on a page/file?

occupied records



- pack them
- keep ptrs to them
- mark start of free space

page header


slot directory

other info (# slots etc)

Variable length records

- Q: How would you store them on a page/file?


occupied records



- how many disk accesses to insert a record?
- to delete one?

page header

Record Formats: Fixed Length




- Information about field types same for all records in a file; stored in *system catalogs*.
- Finding i^{th} field done via arithmetic.

Variable Length records

- Two alternative formats (# fields is fixed):



Fields Delimited by Special Symbols



Array of Field Offsets


Pros and cons?

Variable Length records

- Two alternative formats (# fields is fixed):



Fields Delimited by Special Symbols





Array of Field Offsets

Offset approach: usually superior (direct access to i-th field)

Till now

- Memory hierarchy
- Disks: (>1000x slower) - thus
 - pack info in blocks
 - try to fetch nearby blocks (sequentially)
- Buffer management: very important
 - LRU, MRU, Clock, etc
- Record organization: Slotted page




Hashing

Problem: “*find STU record with ssn=123*”

What if disk space was free, and time was at premium?

Hashing

A: Brilliant idea: key-to-address transformation:




Hashing

- use M , instead of 999,999,999 slots
- hash function: $h(key) = slot-id$

Typically: each hash bucket is a page, holding many records:

STU file



Design decisions - functions

- Goal: uniform spread of keys over hash buckets
- Popular choices:
 - Division hashing
 - $h(x) = (a * x + b) \bmod M$
 - Multiplication hashing
 - $h(x) = \lfloor \text{fractional-part-of} (x * \phi) \rfloor * M$



Hash Design decisions

-functions

-size

-Collision resolution

123; Smith; Main str.




#0 page

#h(123)

M


Collision resolution

linear probing:



Collision resolution


re-hashing



Collision resolution

separate chaining

123; Smith; Main str.



Design decisions - conclusions

- function: division hashing
 - $h(x) = (a * x + b) \bmod M$
- size M : ~90% util.; prime number.
- collision resolution: separate chaining
 - easier to implement (deletions!);
 - no danger of becoming full

Problem with static hashing

- problem: overflow?
- problem: underflow? (underutilization)


Solution: Dynamic/extendible hashing

- idea: shrink / expand hash table on demand..
- ..dynamic hashing

Details: how to grow gracefully, on overflow?

Many solutions - One of them: 'extendible hashing'
[Fagin et al]

Extendible hashing





Extendible hashing


**solution: don't overflow –
instead:**

SPLIT the bucket in two


123; Smith; Main str.




Extendible hashing




Extendible hashing




Extendible hashing



Extendible hashing




Extendible hashing




Extendible hashing

BEFORE



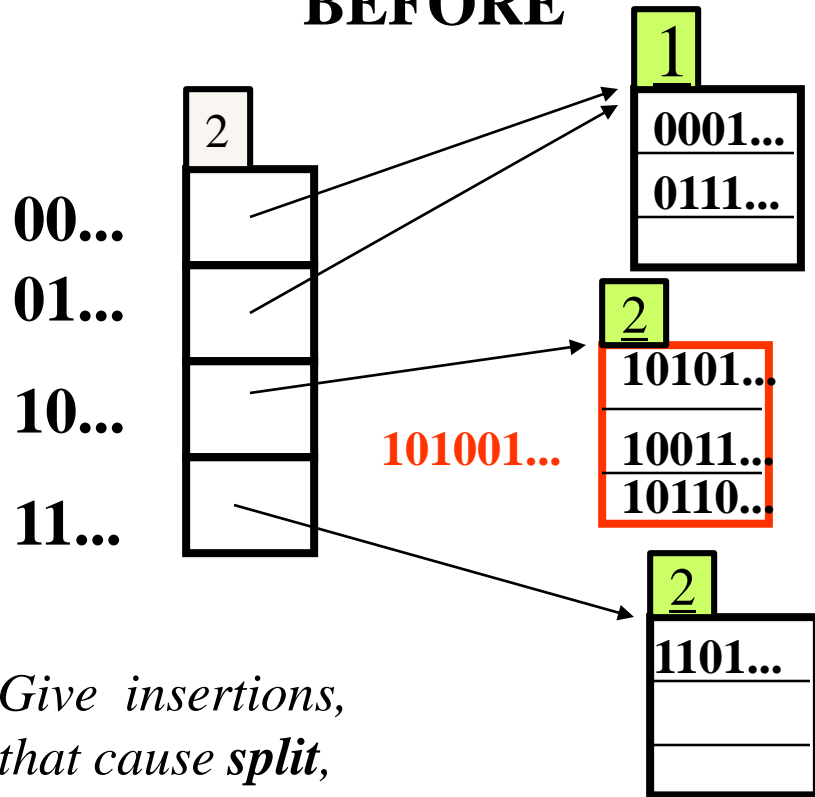
AFTER



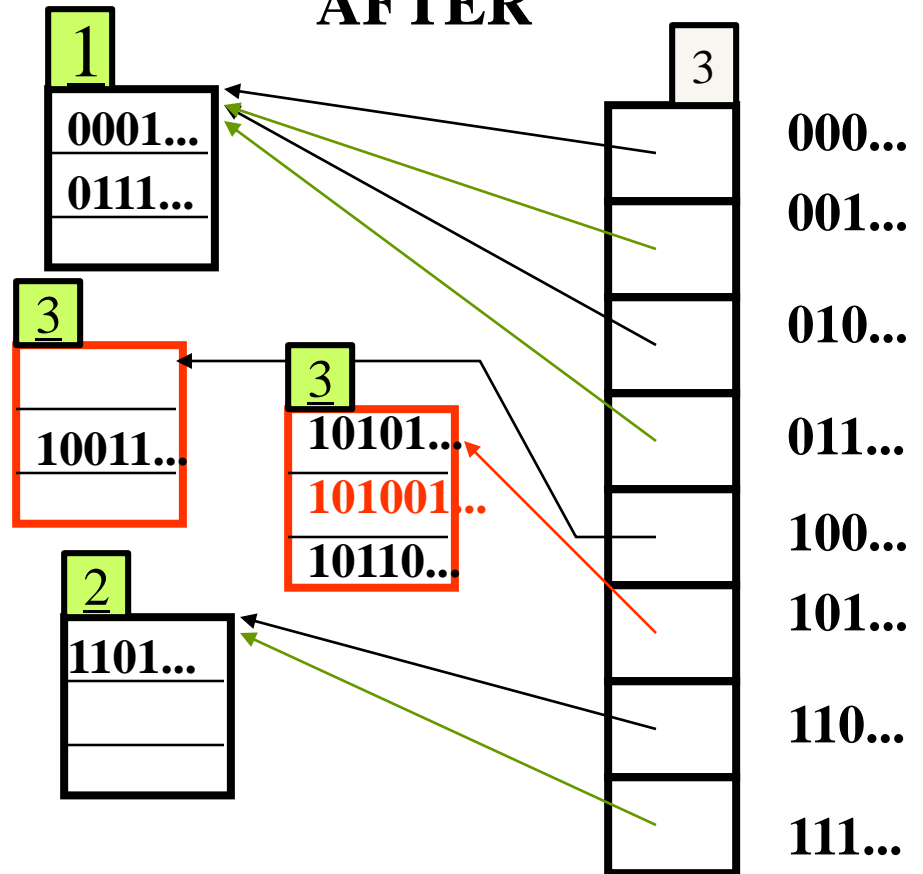
*Do all splits lead to
directory doubling?
A:NO! (most don't)*

Extendible hashing

BEFORE



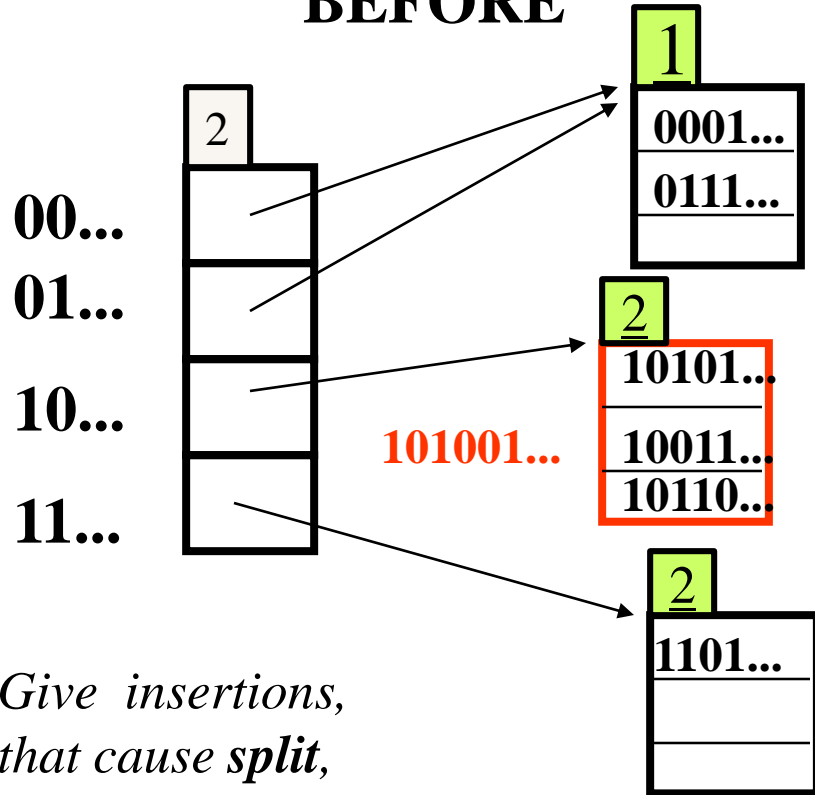
AFTER



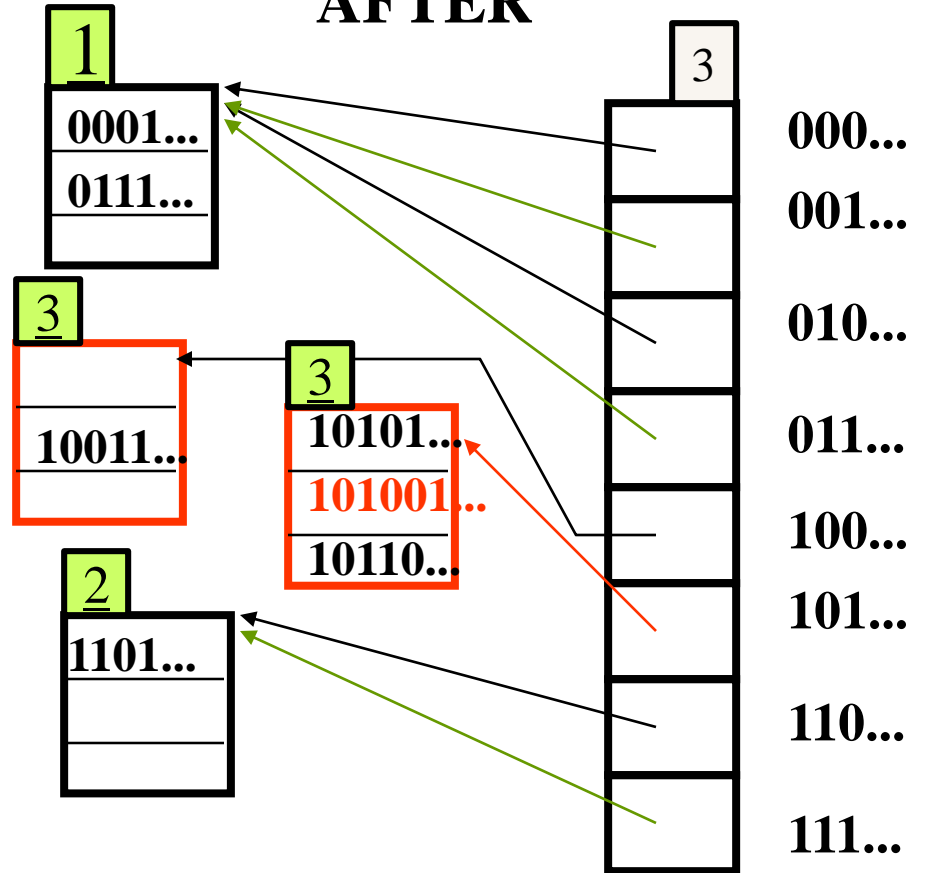
*Give insertions,
that cause **split**,
but **no** directory dup.?*

Extendible hashing

BEFORE



AFTER



Give insertions,
that cause *split*,
but **no** directory dup.?
A: 000... and 001...

Linear hashing

Motivation: can we do something simpler, with smoother growth?

A: split buckets from left to right, **regardless** of which one overflowed ('crazy', but it works well!) -
Eg.:

Linear hashing

Initially: $h(x) = x \bmod N$ (N=4 here)


Assume capacity: 3 records / bucket

Insert key '17'

bucket- id	0	1	2	3
	4 8	5 9 13	6	7 11

Linear hashing

Initially: $h(x) = x \bmod N$ (N=4 here)



Linear hashing

Initially: $h(x) = x \bmod N$ (N=4 here)

overflow of bucket#1

Split #0, anyway!!!

bucket- id	0	1	2	3
	4 8	5 9 13	6	7 11

Linear hashing

Initially: $h(x) = x \bmod N$ (N=4 here)

Split #0, anyway!!!

Q: But, how?

bucket- id	0	1	2	3
	4 8	5 9 13	6	7 11

Linear hashing

A: use two h.f.: $h0(x) = x \bmod N$

$$h1(x) = x \bmod (2*N)$$

17


bucket- id

0

1

2

3



4 8	5 9 13	6	7 11
-------	-------------	---	--------

Linear hashing - after split:

A: use two h.f.: $h0(x) = x \bmod N$

$$h1(x) = x \bmod (2*N)$$

bucket- id	0	1	2	3	4
	8	5 9 13	6	7 11	4


17

Linear hashing - after split:

A: use two h.f.: $h0(x) = x \bmod N$

$$h1(x) = x \bmod (2*N)$$


bucket- id



Linear hashing - after split:


A: use two h.f.: $h0(x) = x \bmod N$

$$h1(x) = x \bmod (2*N)$$



Linear hashing - searching?

- $h0(x) = x \bmod N$ (for the un-split buckets)
- $h1(x) = x \bmod (2*N)$ (for the splitted ones)




Linear hashing - searching?

Q1: find key '6' ?

Q2: find key '4' ?


Q3: key '8' ?



Linear hashing - insertion?

Algo: insert key ' k '

- compute appropriate bucket ' b '
- if the **overflow criterion** is true
 - split the bucket of 'split-ptr'
 - split-ptr ++ (*)




what if we reach the right edge??

Linear hashing - split now?

$h0(x) = x \bmod N$ (for the un-split buckets)

$h1(x) = x \bmod (2*N)$ for the splitted ones)




Linear hashing - split now?

$h0(x) = x \bmod N$ (for the un-split buckets)

$h1(x) = x \bmod (2*N)$ (for the splitted ones)

split ptr



Linear hashing - split now?

~~$h0(x) = x \bmod N$ (for the un-split buckets)~~

$h1(x) = x \bmod (2*N)$ (for the splitted ones)

split ptr



0

1

2


3

4

5

6

7



Linear hashing - split now?

~~$h0(x) = x \bmod N$ (for the un-split buckets)~~

$h1(x) = x \bmod (2*N)$ (for the splitted ones)

split ptr



0

1

2


3

4

5

6

7



Linear hashing - split now?

this state is called **'full expansion'**

split ptr



0

1

2


3

4

5

6

7



Linear hashing - observations

In general, at any point of time, we have at **most two** h.f. active, of the form:

- $h_n(x) = x \bmod (N * 2^n)$

- $h_{n+1}(x) = x \bmod (N * 2^{n+1})$

(after a full expansion, we have only one h.f.)

Hashing - pros?