

Mid-term Summary

- ~ 70% AB
 - How we are going to evaluate in the final?
- Assignment
 - How weekly? – impact: ~88% vs. ~34%
- Feedback
 - Assignment results, separate, acoustic, practical, materials in advance
- Logistics: lecture vs. seminar
- Cheating

Indexing

The background of the slide is a complex, abstract composition. It features a series of overlapping, curved lines in shades of blue, green, and yellow, creating a sense of depth and movement. In the foreground, there are several thick, vibrant lines in magenta, lime green, and dark purple, which appear to be layered over the background. A thin, orange line with small square markers follows a curved path across the middle of the image, adding a technical or data-related element to the design.

Why Indexes?

- Sometimes, we want to retrieve records by specifying the *values in one or more fields*, e.g.,
 - Find all students in the “CS” department
 - Find all students with a gpa > 3
- An index on a file speeds up selections on the *search key fields* for the index (accelerate search on keys)
 - Any subset of the fields of a relation can be the search key for an index on the relation.
 - *Search key* is *not* the same as *key* (e.g., doesn't have to be unique).

Thus, we have:

Index

Data

- ‘data entries’ == what we store at the bottom of the index pages
- what would you use as data entries?
- (3 alternatives here)

Alternatives for Data Entry \mathbf{k}^* in Index

1. Actual data record (with key value \mathbf{k})

123	Smith; Main str; 412-999.9999
-----	-------------------------------

2. $\langle \mathbf{k}, \text{rid of matching data record} \rangle$

\$40	Rid-1
------	-------

\$40	Rid-2
------	-------

...

3. $\langle \mathbf{k}, \text{list of rids of matching data records} \rangle$

\$40	Rid-1	Rid-2	...
------	-------	-------	-----

Alternatives for Data Entry k^* in Index

1. Actual data record (with key value k)
 2. $\langle k, \text{rid of matching data record} \rangle$
 3. $\langle k, \text{list of rids of matching data records} \rangle$
- Choice is orthogonal to the indexing technique.
 - Examples of indexing techniques: B+ trees, hash-based structures, R trees, ...
 - Typically, index contains auxiliary info that directs searches to the desired data entries
 - Can have multiple (different) indexes per file.
 - E.g. file sorted on *age*, with a hash index on *name* and a B-tree index on *salary*.

Alternatives for Data Entries (Contd.)

Alternative 1:

Actual data record (with key value **k**)

- Then, this is a clustering/sparse index, and constitutes a file organization (like Heap files or sorted files).
- **At most one** index on a given collection of data records can use Alternative 1.
- Saves pointer lookups but can be expensive to maintain with insertions and deletions.

Alternatives for Data Entries (Contd.)

Alternative 2

<k, rid of matching data record>

and Alternative 3

<k, list of rids of matching data records>

- Easier to maintain than Alternative 1.
- If more than one index is required on a given file, at most one index can use Alternative 1; rest must use Alternatives 2 or 3.
- Alternative 3 more compact than Alternative 2, but leads to *variable sized data entries* even if search keys are of fixed length.
- Even worse, for large rid lists the data entry would have to span multiple pages!

Outline – Tree Structured Index

- Motivation
- ISAM
- B-trees
- Tree vs. Hash-based index
- Index organization: clustered vs. nonclustered

Motivation

- Accelerate searches on big indexes
- How to support range searches?
- equality searches?

Range Searches

- “*Find all students with $gpa > 3.0$* ”
- may be slow, even on sorted file
- What to do?



Range Searches

- “*Find all students with $gpa > 3.0$* ”
- may be slow, even on sorted file
- Solution: Create an ‘index’ file.




Range Searches

- More details:
- if index file is small, do binary search there
- Otherwise??




ISAM

- Repeat recursively!



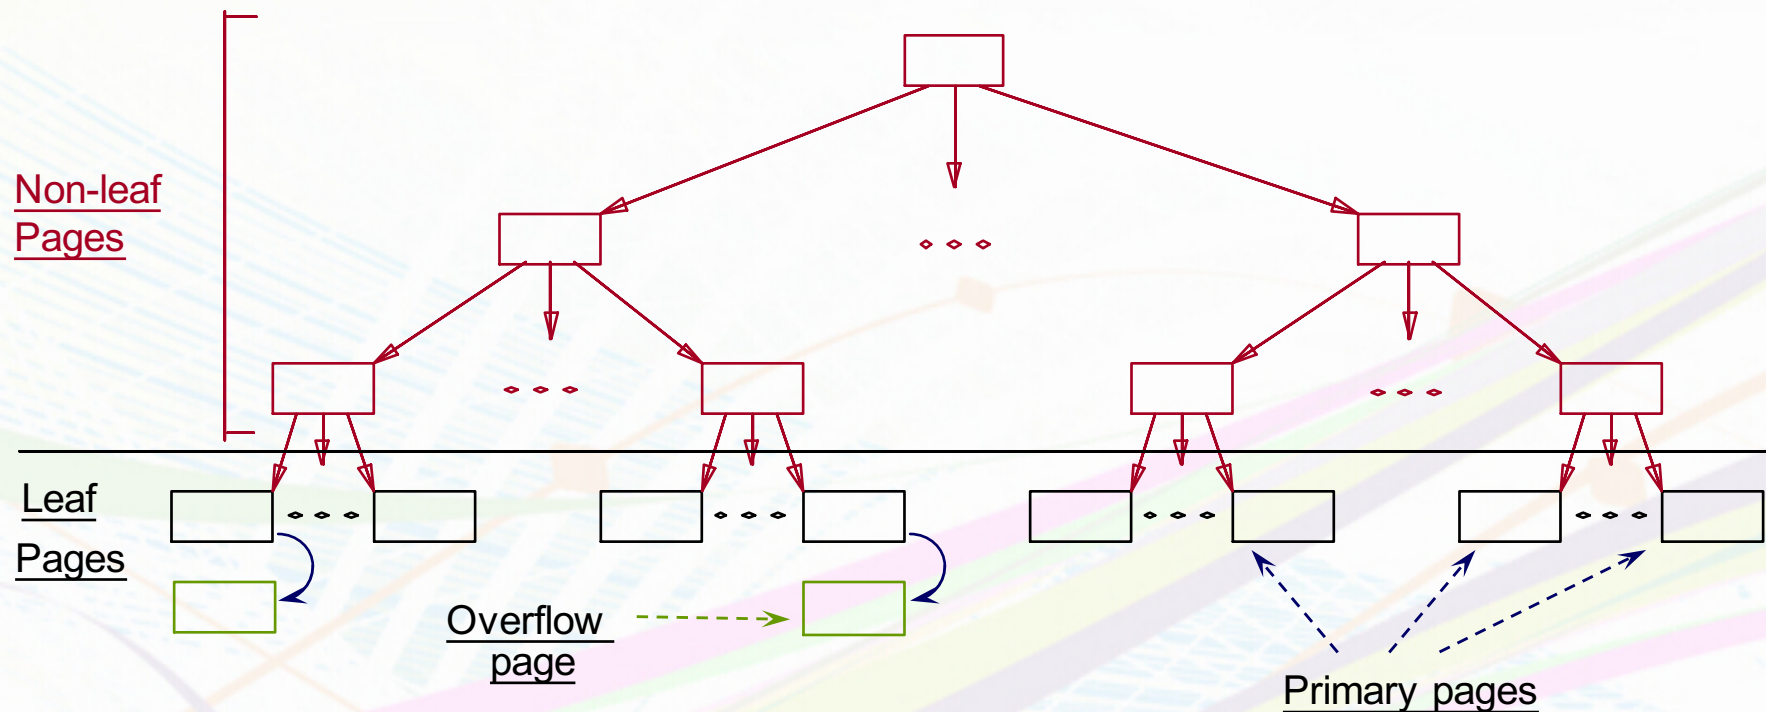
ISAM

- OK - what if there are insertions and overflows?




ISAM

- Overflow pages, linked to the primary page




Example ISAM Tree

- 2 entries per page




ISAM



Details

- format of an index page?
- how full should a newly created ISAM be?

ISAM




Details

- format of an index page?
- how full should a newly created ISAM be?
 - ~80-90% (not 100%)


ISAM is a STATIC Structure

- that is, index pages don't change
- *File creation:* Leaf (data) pages allocated sequentially, sorted by search key; then index pages allocated, then overflow pgs.




ISAM is a STATIC Structure

- *Search:* Start at root; use key comparisons to go to leaf.
- Cost = ??




ISAM is a STATIC Structure

- *Search:* Start at root; use key comparisons to go to leaf.
- Cost = ??



ISAM is a STATIC Structure


- Search: Start at root; use key comparisons to go to leaf.
- Cost = $\log_F N$;
- $F = \# \text{ entries/pg (i.e., fanout),}$
- $N = \# \text{ leaf pgs}$



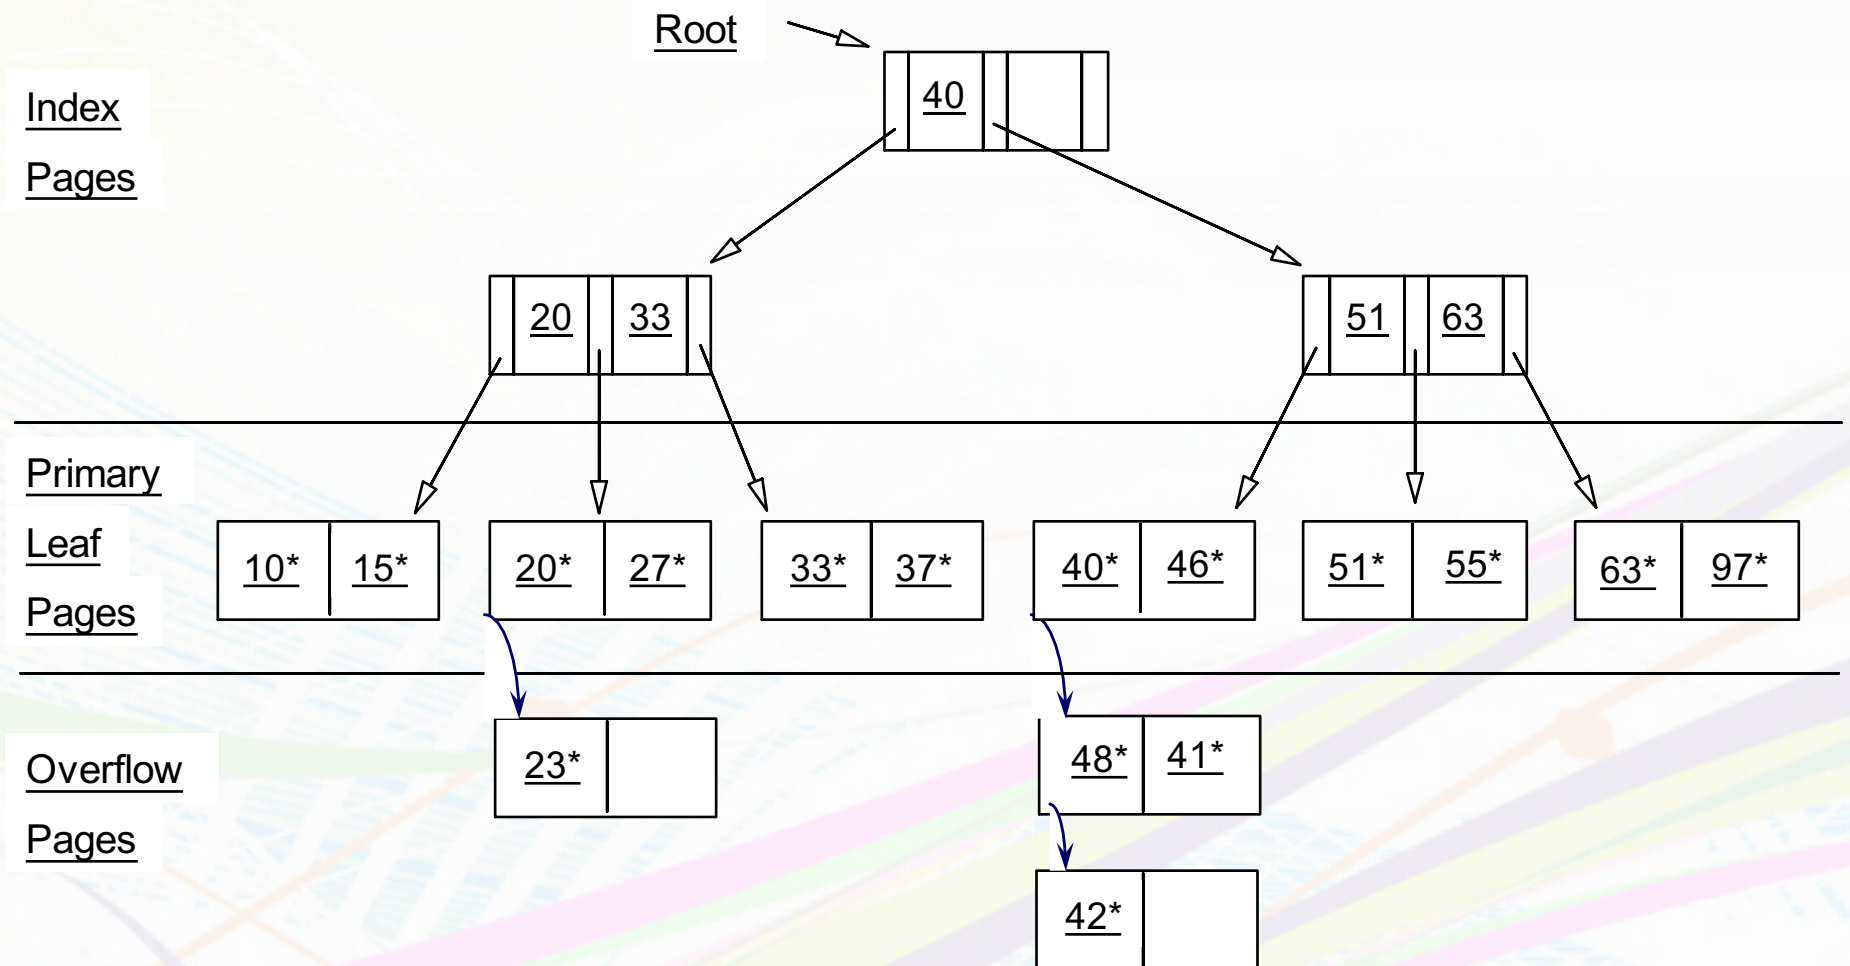
ISAM is a STATIC Structure

Insert: Find leaf that data entry belongs to, and put it there. Overflow page if necessary.

Delete: Find and remove from leaf; if empty page, de-allocate.




Example: Insert 48*, 23*, 41*, 42*

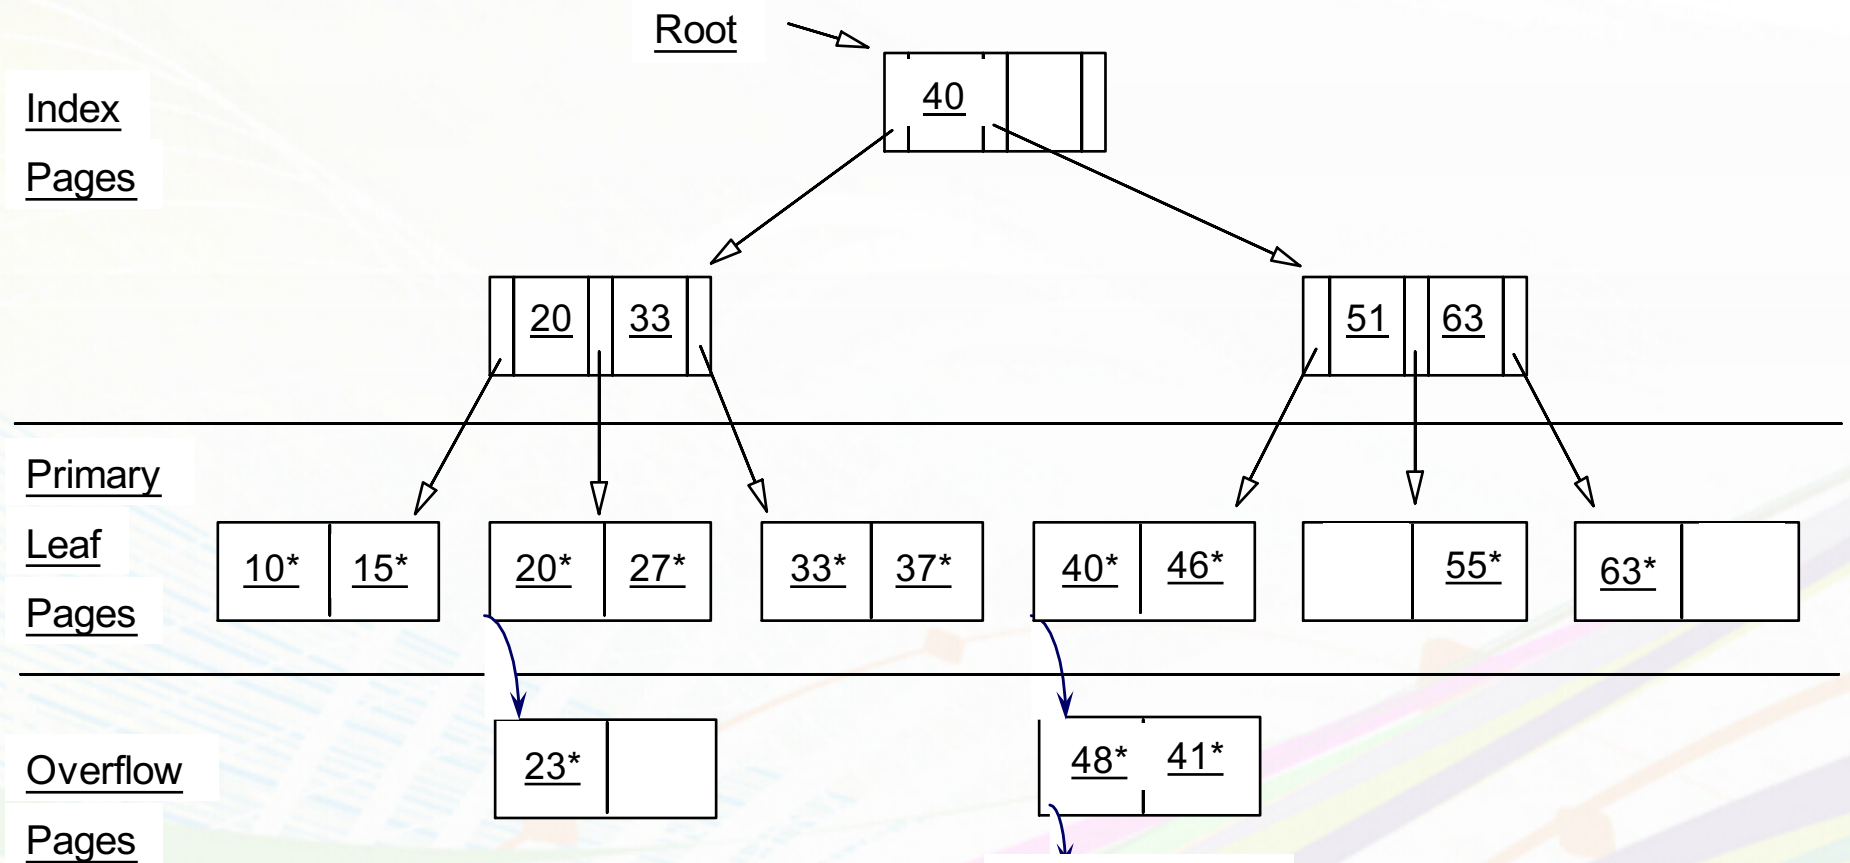


21* means

- <21> + the rest of the record
- (it's a bit more complicated – but we stay with that, for the moment).
- '21' plain means just 4 bytes, to store integer 21



... then delete 42*, 51*, 97*



☞ Note that 51* appears in index levels, but not in leaf!

ISAM ---- Issues?

- Pros
 - ????
- Cons
 - ????

Outline

- Motivation
- ISAM
- B-trees
- Tree vs. Hash-based index
- Index organization: clustered vs. nonclustered

B-trees and Why?

- the most successful family of index schemes (B-trees, B⁺-trees, B^{*}-trees)
- Can be used for primary/secondary, clustering/non-clustering index.
- balanced “n-way” search trees
- Flexible and dynamic


B-trees

[Rudolf Bayer and McCreight, E. M.
Organization and Maintenance of Large
Ordered Indexes. Acta Informatica 1,
173-189, 1972.]




B-trees Example

B-tree of order $d=1$:



B - tree properties:

- each node, in a B-tree of order d :
 - Key order
 - at most $n=2d$ keys
 - at least d keys (except root, which may have just 1 key)
 - all leaves at the same level
 - if number of pointers is k , then node has exactly $k-1$ keys
 - (leaves are with data entries)




Properties

- “block aware” nodes: each node -> disk page
- $O(\log(N))$ for everything! (ins/del/search)
- typically, if $d = 50 - 100$, then 2 - 3 levels
- utilization $\geq 50\%$, guaranteed; on average 69%

Queries

- Algo for exact match query? (eg., ssn=8?)




JAVA animation!

<http://slady.net/java/bt/>

strongly recommended! (with all usual precautions – VM etc)


Queries

- Algo for exact match query? (eg., ssn=8?)




Queries

- Algo for exact match query? (eg., ssn=8?)




Queries

- Algo for exact match query? (eg., ssn=8?)



Queries

- Algo for exact match query? (eg., ssn=8?)




Queries

- what about range queries? (eg., $5 < salary < 8$)
- Proximity/ nearest neighbor searches? (eg., $salary \sim 8$)


Queries

- what about range queries? (eg., $5 < \text{salary} < 8$)
- Proximity/ nearest neighbor searches? (eg., $\text{salary} \sim 8$)




Queries

- what about range queries? (eg., $5 < salary < 8$)
- Proximity/ nearest neighbor searches? (eg., $salary \sim 8$)




Queries

- what about range queries? (eg., $5 < \text{salary} < 8$)
- Proximity/ nearest neighbor searches? (eg., *salary ~ 8*)



Queries

- what about range queries? (eg., $5 < \text{salary} < 8$)
- Proximity/ nearest neighbor searches? (eg., *salary* ~ 8)




B-trees: Insertion

- Insert in leaf; on overflow, push middle up (recursively)
- split: preserves B - tree properties


B-trees

Easy case: Tree T0; insert '8'




B-trees

Tree T0; insert '8'



B-trees


Hardest case: Tree T0; insert '2'



2


B-trees

Hardest case: Tree T0; insert '2'



B-trees


Hardest case: Tree T0; insert '2'



B-trees

Hardest case: Tree T0; insert '2'

Final state



B-trees: Insertion

- Insert in leaf; on overflow, push middle up (recursively – ‘propagate split’)
- split: preserves all B - tree properties (!!)
- notice how it grows: height increases when root overflows & splits
- Automatic, incremental re-organization (contrast with ISAM!)

Pseudo-code

INSERTION OF KEY 'K'

find the correct leaf node 'L' ;

if ('L' overflows){

 split 'L' , and push middle key to parent node 'P' ;

 if ('P' overflows){

 repeat the split recursively; }

else{

 add the key 'K' in node 'L' ;

 /* maintaining the key order in 'L' */ }

Overview

- ...
- B – trees
 - Dfn, Search, insertion, deletion
- ...

Deletion


Rough outline of algo:

- Delete key;
- on underflow, may need to merge

In practice, some implementors just allow underflows to happen...


B-trees – Deletion

Easiest case: Tree T0; delete '3'



B-trees – Deletion

Easiest case: Tree T0; delete '3'




B-trees – Deletion

- Case1: delete a key at a leaf – no underflow
- Case2: delete non-leaf key – no underflow
- Case3: delete leaf-key; underflow, and ‘rich sibling’
- Case4: delete leaf-key; underflow, and ‘poor sibling’


B-trees – Deletion

- Case1: delete a key at a leaf – no underflow (delete 3 from T0)




B-trees – Deletion

- Case2: delete a key at a non-leaf – no underflow (eg., delete 6 from T0)




B-trees – Deletion

- Case2: delete a key at a non-leaf – no underflow (eg., delete 6 from T0)



B-trees – Deletion


- Case2: delete a key at a non-leaf – no underflow (eg., delete 6 from T0)



B-trees – Deletion

- Case2: delete a key at a non-leaf – no underflow (eg., delete 6 from T0)

FINAL TREE



B-trees – Deletion


- Case2: delete a key at a non-leaf – no underflow (eg., delete 6 from T0)
- Q: How to promote?
- A: pick the largest key from the left sub-tree (or the smallest from the right sub-tree)
- Observation: every deletion eventually becomes a deletion of a leaf key

B-trees – Deletion

- Case1: delete a key at a leaf – no underflow
- ⇒ • Case2: delete non-leaf key – no underflow
- Case3: delete leaf-key; underflow, and ‘rich sibling’
- Case4: delete leaf-key; underflow, and ‘poor sibling’


B-trees – Deletion

- Case3: underflow & ‘rich sibling’ (eg., delete **7** from T0)



B-trees – Deletion

- Case3: underflow & ‘rich sibling’ (eg., delete **7** from T0)

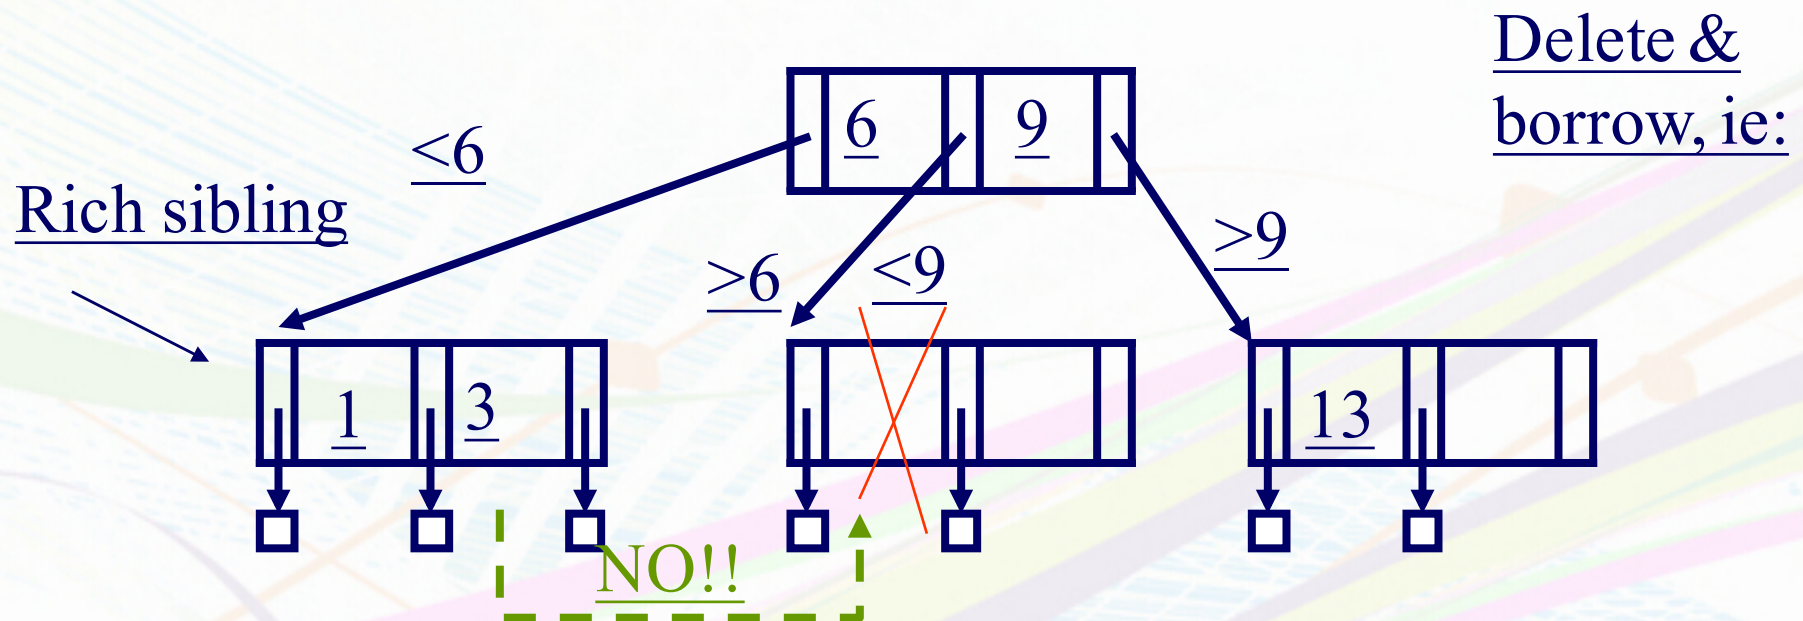


B-trees – Deletion

- Case3: underflow & ‘rich sibling’
- ‘rich’ = can give a key, without underflowing
- ‘borrowing’ a key: THROUGH the PARENT!


B-trees – Deletion

- Case3: underflow & ‘rich sibling’ (eg., delete 7 from T0)




B-trees – Deletion

- Case3: underflow & ‘rich sibling’ (eg., delete 7 from T0)



B-trees – Deletion


- Case3: underflow & ‘rich sibling’ (eg., delete **7** from T0)



B-trees – Deletion

- Case3: underflow & ‘rich sibling’ (eg., delete **7** from T0)

FINAL TREE




B-trees – Deletion

- Case1: delete a key at a leaf – no underflow
- Case2: delete non-leaf key – no underflow
- Case3: delete leaf-key; underflow, and
⇒ ‘rich sibling’
- Case4: delete leaf-key; underflow, and
‘poor sibling’


B-trees – Deletion

- Case4: underflow & ‘poor sibling’ (eg., delete 13 from T0)




B-trees – Deletion

- Case4: underflow & ‘poor sibling’ (eg., delete 13 from T0)



B-trees – Deletion

- Case4: underflow & ‘poor sibling’ (eg., delete 13 from T0)




B-trees – Deletion

- Case4: underflow & ‘poor sibling’ (eg., delete 13 from T0)
- Merge, by pulling a key from the parent
- exact reversal from insertion: ‘split and push up’, vs. ‘merge and pull down’
- Ie.:

B-trees – Deletion


- Case4: underflow & ‘poor sibling’ (eg., delete 13 from T0)



B-trees – Deletion

- Case4: underflow & ‘poor sibling’ (eg., delete 13 from T0)

FINAL TREE



B-trees – Deletion

- Case4: underflow & ‘poor sibling’
- -> ‘pull key from parent, and merge’
- Q: What if the parent underflows?
- A: repeat recursively

B-tree deletion - pseudocode

DELETION OF KEY 'K'

locate key 'K', in node 'N'

if('N' is a non-leaf node) {

delete 'K' from 'N';

find the immediately largest key 'K1';

/* which is guaranteed to be on a leaf node 'L' */

copy 'K1' in the old position of 'K';

invoke this DELETION routine on 'K1' from the leaf
node 'L';

else {

/* 'N' is a leaf node */

... (next slide..)

B-tree deletion - pseudocode

```
/* ' N' is a leaf node */  
if( ' N' underflows ){  
    let ' N1' be the sibling of ' N' ;  
    if( ' N1' is "rich"){ /* ie., N1 can lend us a key */  
        borrow a key from ' N1' THROUGH the parent node;  
    }else{ /* N1 is 1 key away from underflowing */  
        MERGE: pull the key from the parent ' P',  
        and merge it with the keys of ' N' and ' N1' into a new  
        node;  
        if( ' P' underflows){ repeat recursively }  
    }  
}
```

Outline

- Motivation
- ISAM
- B-trees
- Tree vs. Hash-based index
- Index organization: clustered vs. nonclustered

Tree vs. Hash-based index

- Hash-based index

- Good for equality selections.

- File = a collection of *buckets*. Bucket = *primary page* plus 0 or more *overflow pages*.

- *Hash function h*: $h(r.search_key) =$ bucket in which record r belongs.

- Tree-based index

- Good for range selections, big indexes

- Hierarchical structure (Tree) directs searches

- Leaves contain data entries sorted by search key value

- **B+ tree**: all root->leaf paths have equal length (*height*)

Outline

- Motivation
- ISAM
- B-trees
- Tree vs. Hash-based index
- Index organization: clustered vs. nonclustered

Indexing - clustered index example

Clustering/sparse index on ssn

123
456
...

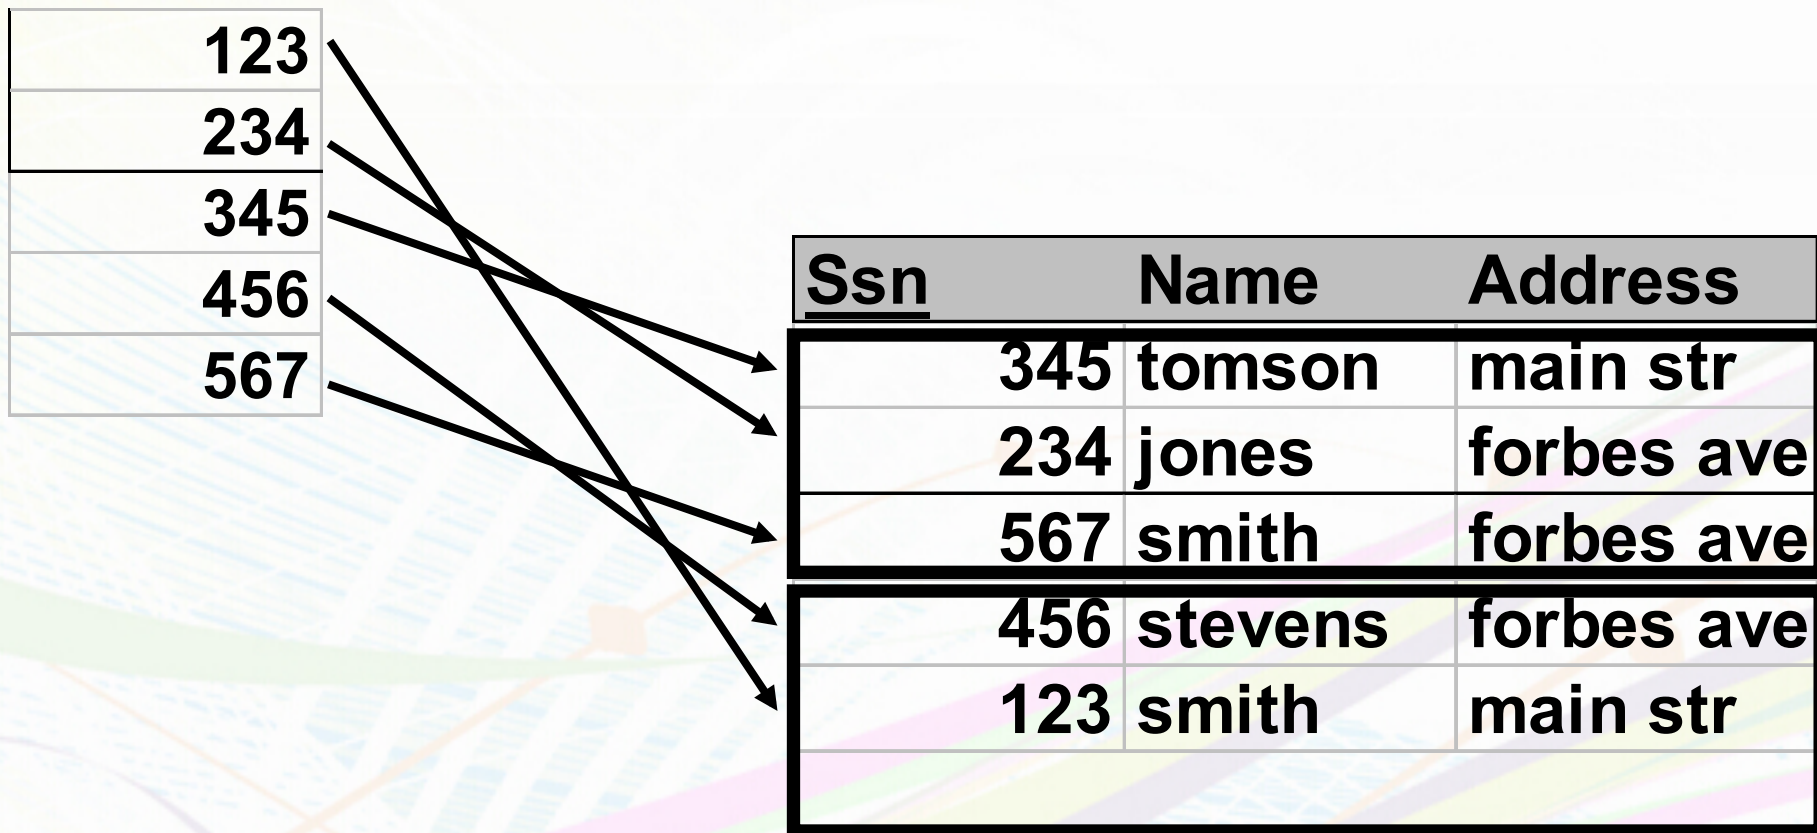
≥ 123

≥ 456

STUDENT			
<u>Ssn</u>	Name	Address	
123	smith	main str	
234	jones	forbes ave	
345	tomson	main str	
456	stevens	forbes ave	
567	smith	forbes ave	

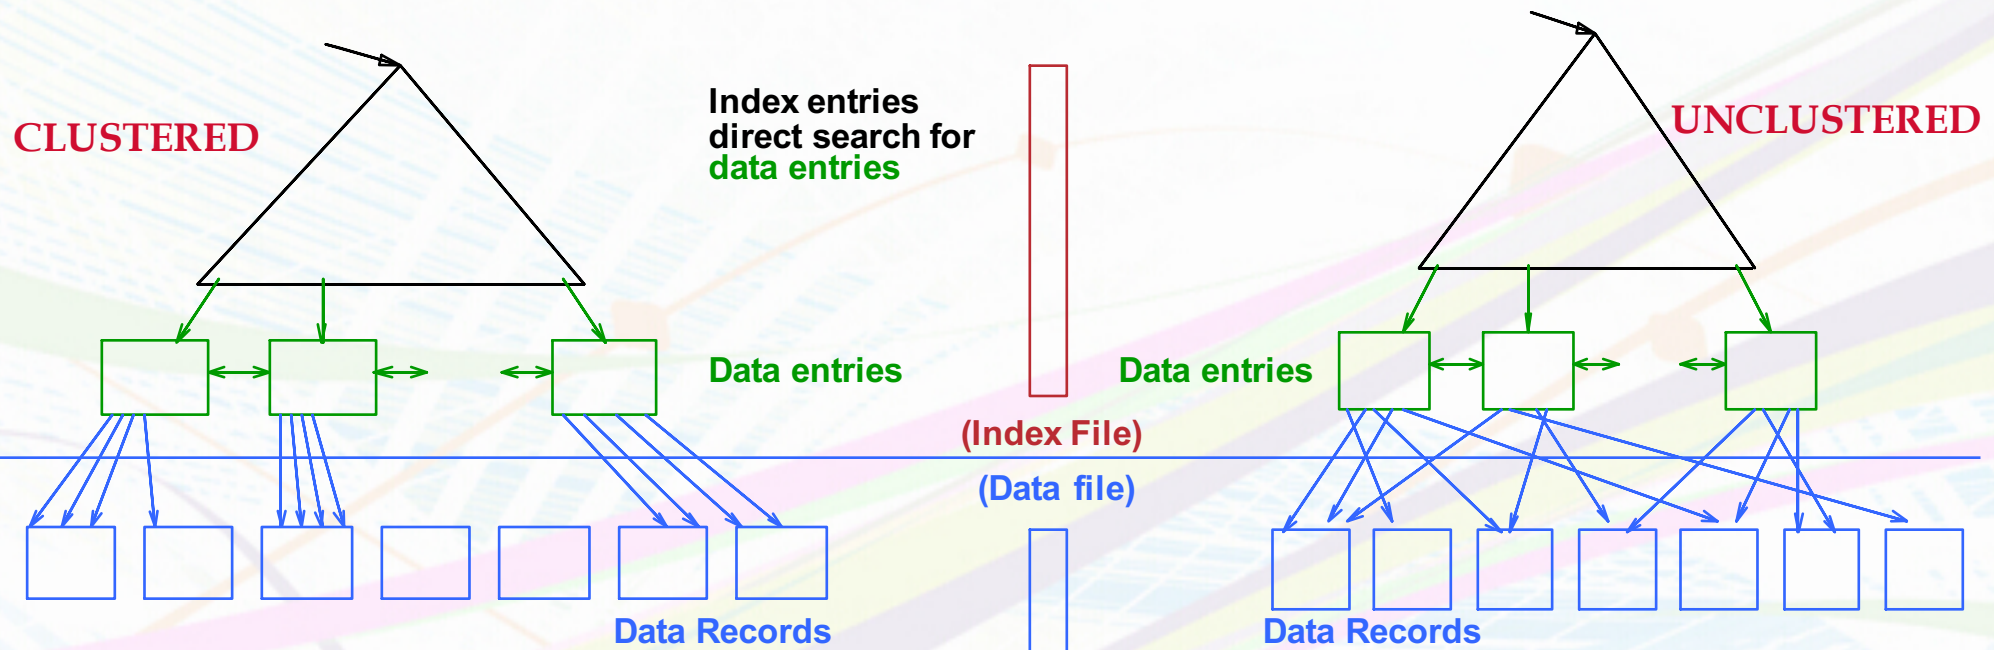
Indexing - non-clustered

Non-clustering / dense index



Index Classification - clustered

- *Clustered vs. unclustered*: If order of **data records** is the same as, or 'close to', order of **index data entries**, then called *clustered index*.



Index Classification - clustered

- A file can have a clustered index on at **most one** search key.
- Cost of retrieving data records through index varies *greatly* based on whether index is clustered!
- Note: Alternative 1 implies clustered, *but not vice-versa*.

But, for simplicity, you may think of them as equivalent..

Clustered vs. Unclustered Index

- Cost of retrieving records found in range scan:
 - Clustered: cost =
 - Unclustered: cost \approx
- What are the tradeoffs????

Clustered vs. Unclustered Index

- Cost of retrieving records found in range scan:
 - Clustered: cost = # pages in file w/matching records
 - Unclustered: cost \approx # of matching index data entries
- What are the tradeoffs????

Clustered vs. Unclustered Index

- Cost of retrieving records found in range scan:
 - Clustered: cost = # pages in file w/matching records
 - Unclustered: cost \approx # of matching index data entries
- What are the tradeoffs????
 - Clustered Pros:
 - Efficient for range searches
 - May be able to do some types of compression
 - Clustered Cons:
 - Expensive to maintain (on the fly or sloppy with reorganization)

Chapter 8, 10

B+ tree

Exercises: 10.2-10.5, 10.8