



SQL

NoSQL



NoSQL: The Name

- “SQL” = Traditional relational DBMS
- Recognition over past decade or so:
Not every data management/analysis problem is best solved using a traditional relational DBMS
- “NoSQL” = “No SQL” =
Not using traditional relational DBMS
- “No SQL” \neq Don't use SQL language



NoSQL: The Name

- “SQL” = Traditional relational DBMS
- Recognition over past decade or so:
Not every data management/analysis problem is best solved using a traditional relational DBMS
- “NoSQL” = “No SQL” =
Not using traditional relational DBMS
- “No SQL” \neq Don't use SQL language
- * “NoSQL” = “Not Only SQL”



Not Only SQL

Not every data management/analysis problem is best solved using a traditional DBMS



**Database Management System (DBMS)
provides....**

... efficient, reliable, convenient, and safe
multi-user storage of and access to massive
amounts of persistent data.



NoSQL Systems

Alternative to traditional relational DBMS

- + Flexible schema
- + Quicker/cheaper to set up
- + Massive scalability
- + Relaxed consistency → higher performance & availability
- No declarative query language → more programming
- Relaxed consistency → fewer guarantees



Example #1: Web log analysis

Each record: UserID, URL, timestamp, additional-info

Task: Find all pairs of UserIDs accessing same URL



Example #1: Web log analysis

Each record: UserID, URL, timestamp, additional-info

Separate records: UserID, name, age, gender, ...

Task: Find average age of user accessing given URL



Example #2: Social-network graph

Each record: UserID₁, UserID₂

Separate records: UserID, name, age, gender, ...

Task: Find all friends of given user



Example #2: Social-network graph

Each record: UserID₁, UserID₂

Separate records: UserID, name, age, gender, ...

Task: Find all friends of friends given user



Example #2: Social-network graph

Each record: UserID₁, UserID₂

Separate records: UserID, name, age, gender, ...

Task: Find all women friends of men friends of given user



Example #2: Social-network graph

Each record: UserID₁, UserID₂

Separate records: UserID, name, age, gender, ...

Task: Find all friends of friends of friends of ... friends of given user



Example #3: Wikipedia pages

Large collection of documents

Combination of structured and unstructured data

Task: Retrieve area of all counties in square meter



NoSQL Systems

Alternative to traditional relational DBMS

- + Flexible schema
- + Quicker/cheaper to set up
- + Massive scalability
- + Relaxed consistency → higher performance & availability
- No declarative query language → more programming
- Relaxed consistency → fewer guarantees

NoSQL Systems

Several realization



NoSQL Systems

Several realization

- MapReduce framework
- Key-value stores
- Document stores
- Graph database systems





MapReduce Framework

Originally from Google, open source Hadoop

- No data model, data stored in files
- User provides specific functions




Map and Reduce Functions

Map: Divide problem into subproblems


Reduce: Do work on subproblems, combine results

- **map:** takes (item_key, value), produces one or more (key, value') pairs
- **reduce:** takes (key, {set of value'}), produces one or more output results
 - typically (key, agg_value)

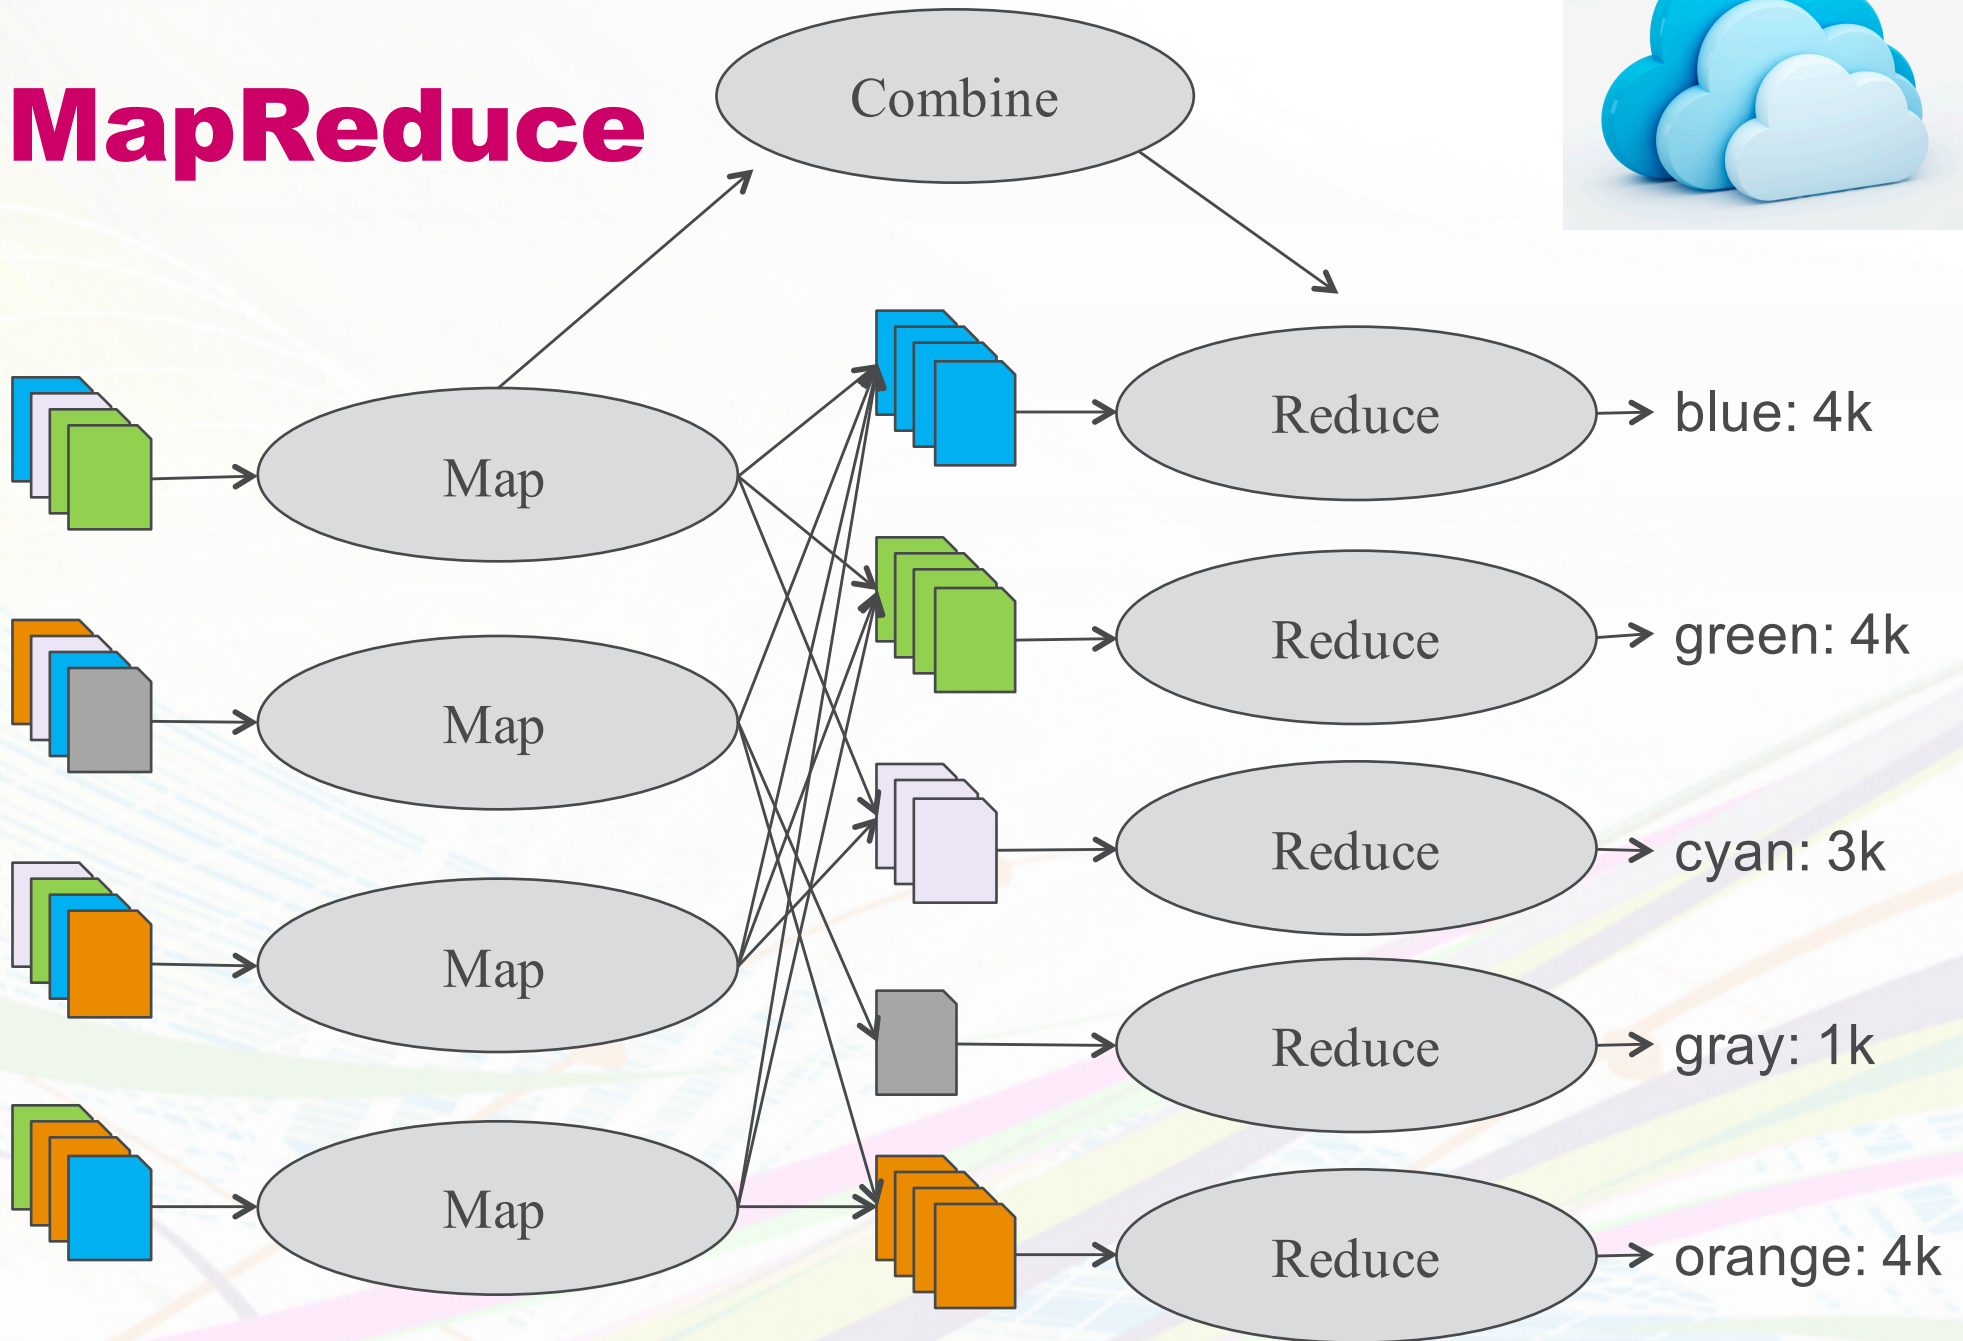
MapReduce



MapReduce

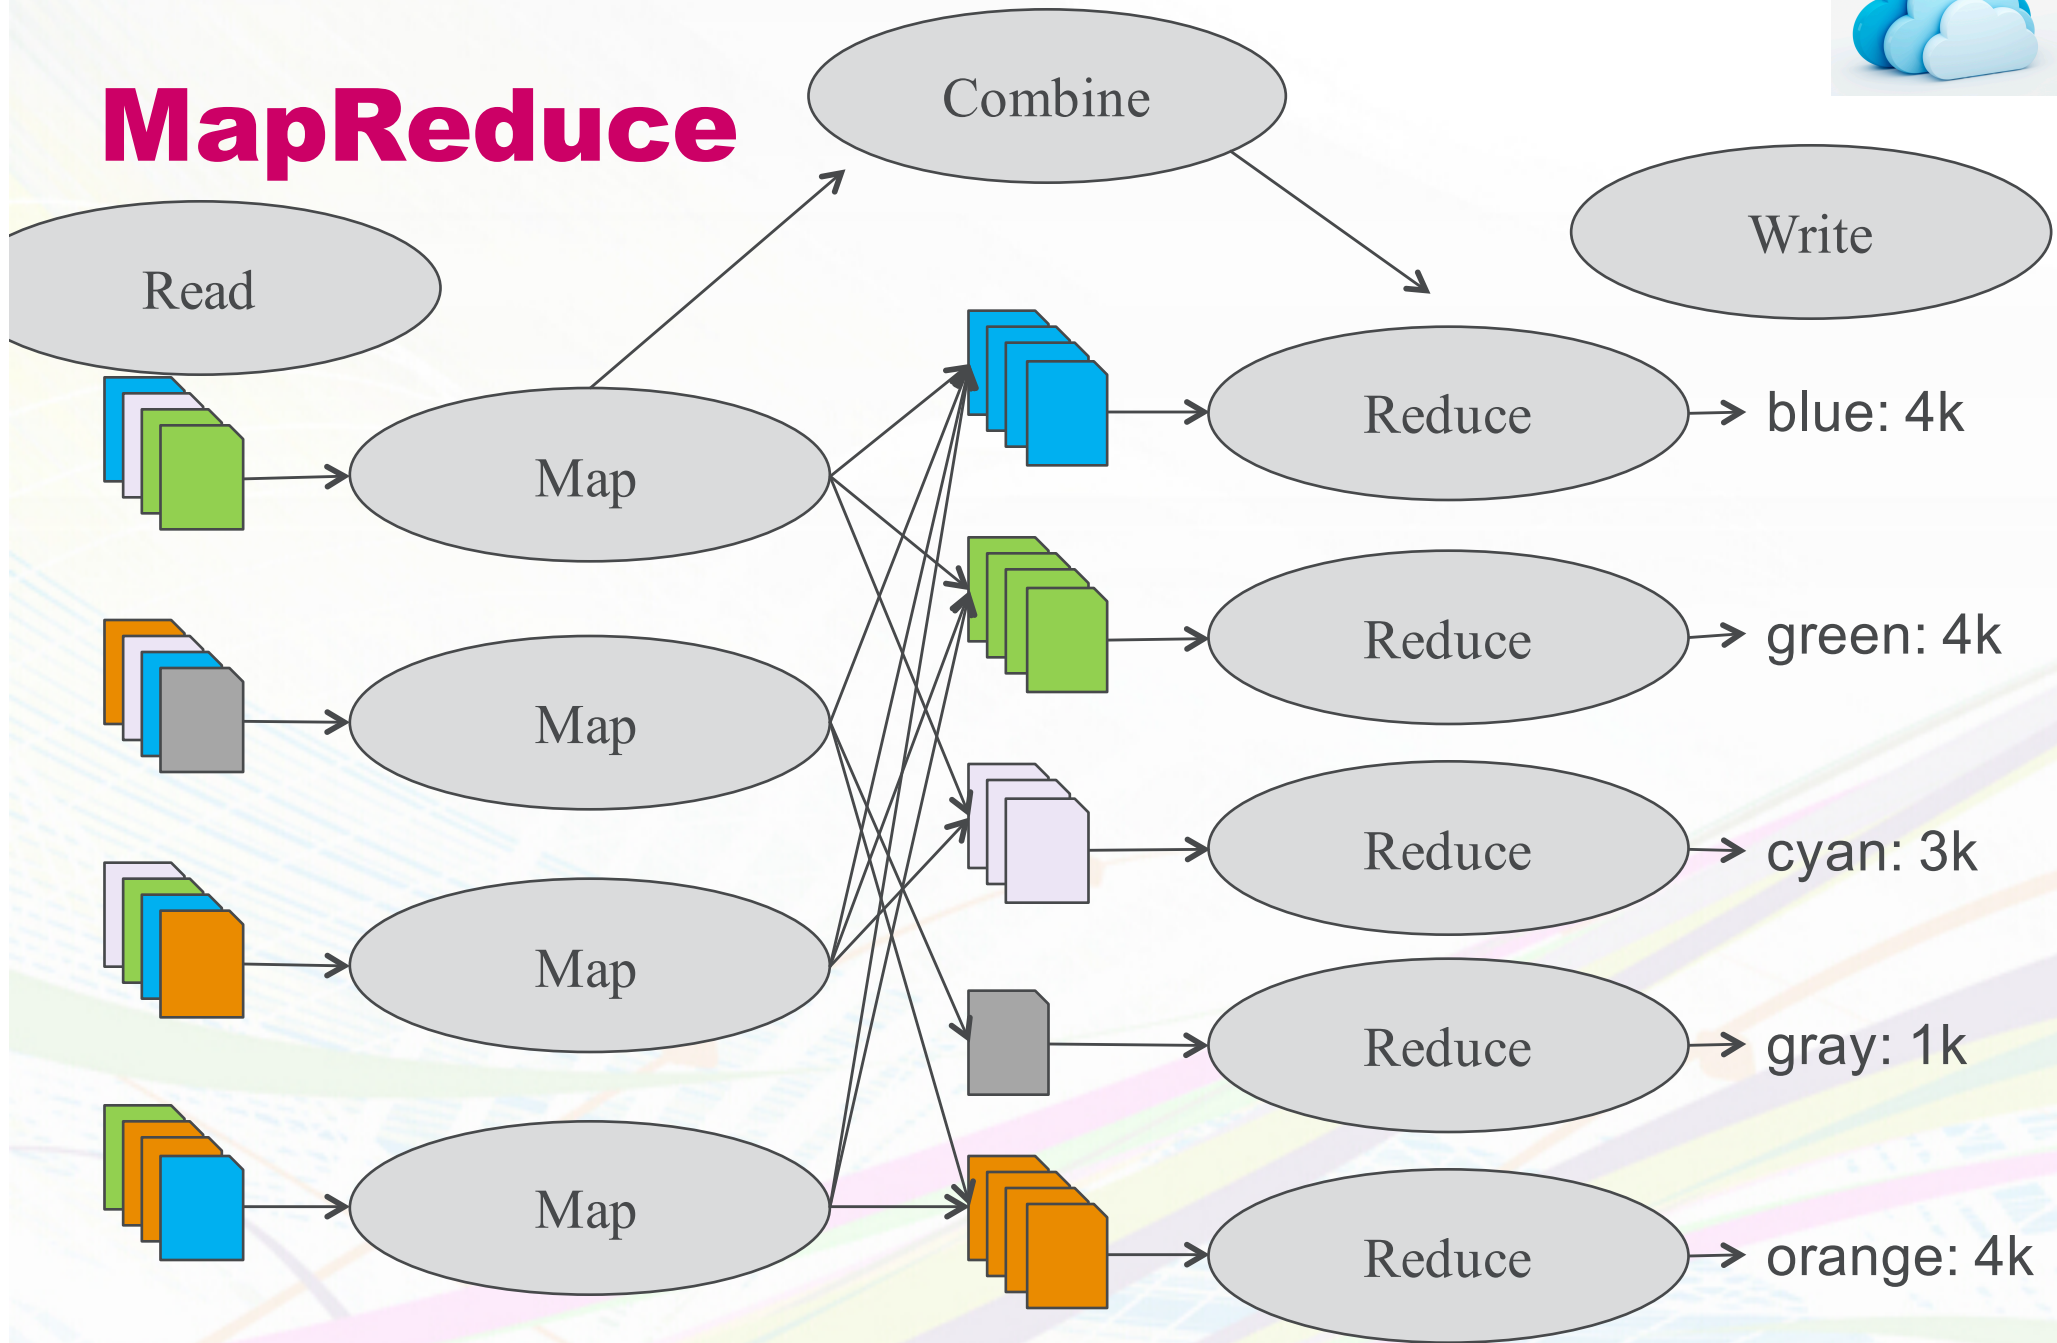


MapReduce





MapReduce



Simple example: Word count



```
map(String key, String value) {  
    // key: document name, line no  
    // value: contents of line  
    for each word w in value:  
        emit(w, "1")  
}
```

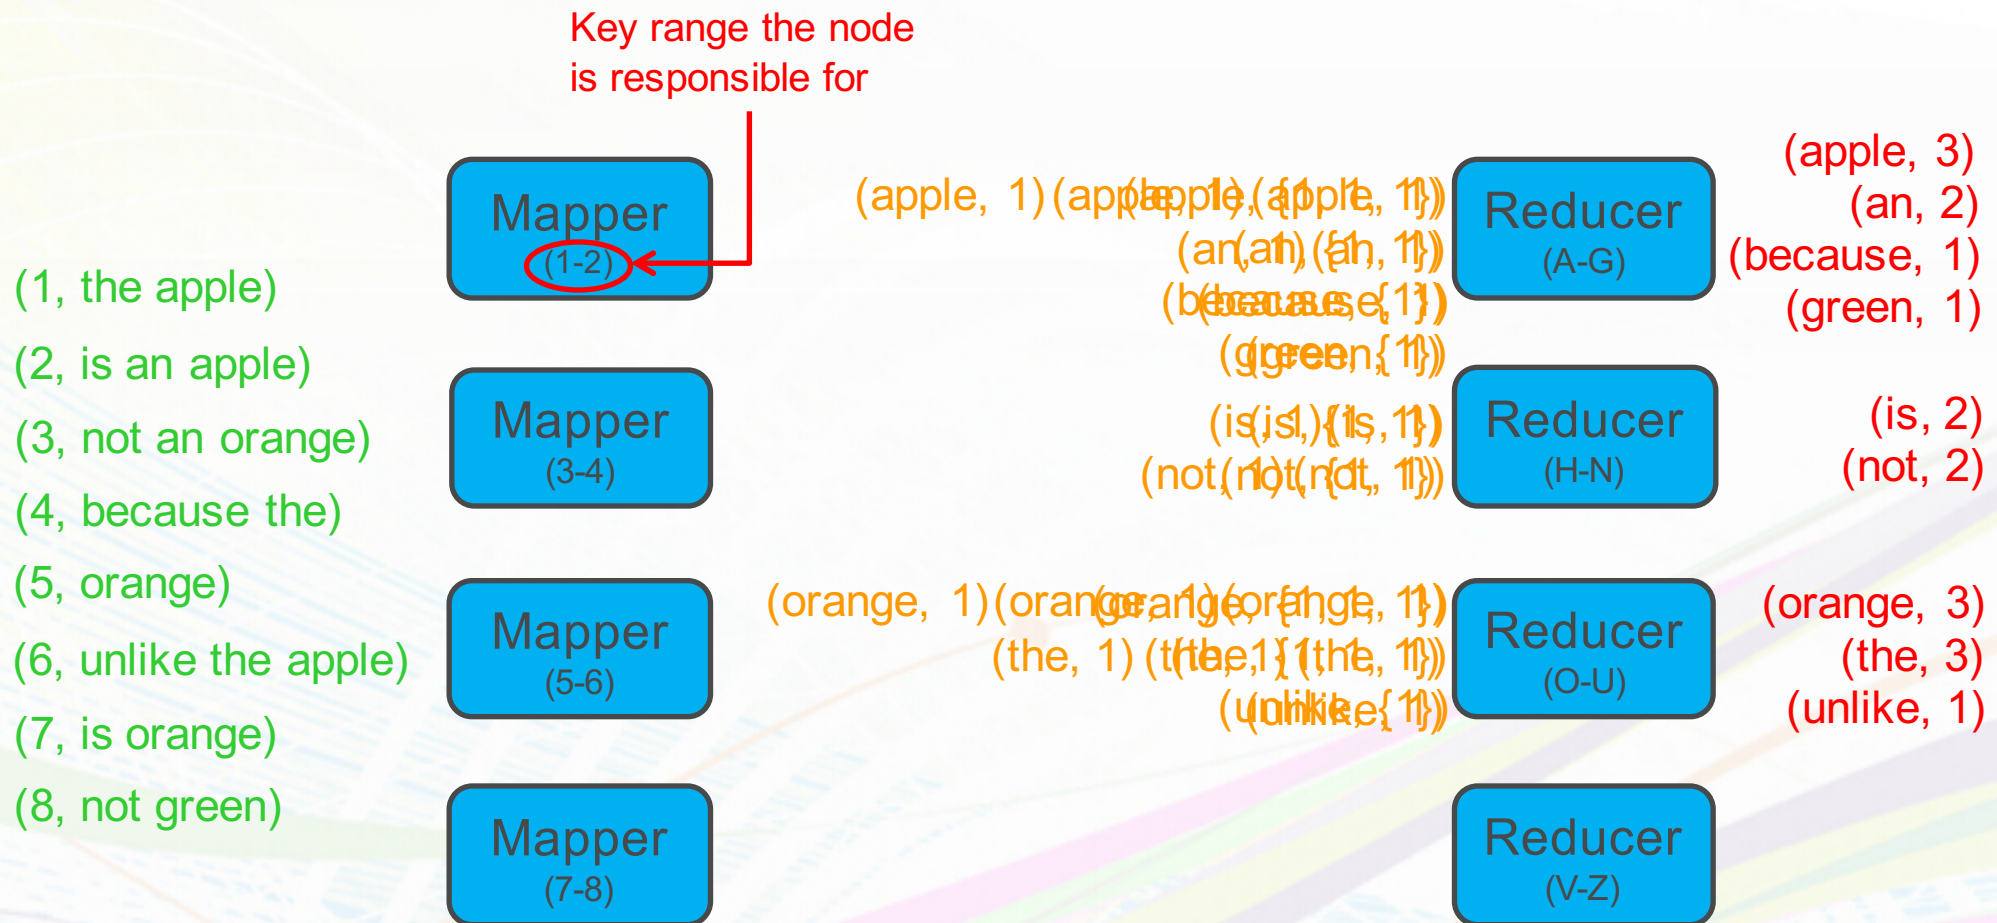
```
reduce(String key, Iterator values) {  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    emit(key, result)  
}
```

Goal: Given a set of documents, count how often each word occurs

Input: Key-value pairs (document:lineNumber, text)

Output: Key-value pairs (word, #occurrences)

Simple example: Word count





MapReduce Framework

- No data model, data stored in files
- User provides specific functions
- System provides data processing “glue”, fault tolerance, scalability



MapReduce Framework

Schemas and declarative queries are missed

Hive – schemas, SQL-like query language

Pig – more imperative but with relational operators

- Both compile to “workflow” of Hadoop (MapReduce) jobs

Dryad allows user to specify workflow

- Also DryadLINQ language



Key-Value Stores

Extremely simple interface

- Data model: (key, value) pairs
- Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)

Implementation: efficiency, scalability, fault-tolerance

- Records distributed to nodes based on key
- Replication
- Single-record transactions, “eventual consistency”



Key-Value Stores

Extremely simple interface

- Data model: (key, value) pairs
- Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)

Example systems

- Google BigTable, Amazon Dynamo, Cassandra, HBase, ...



Document Stores

Like Key-Value Stores except value is document

- **Data model:** (key, document) **pairs**
- **Document:** JSON, XML, other semistructured formats
- **Basic operations:** Insert(key,document), Fetch(key), Update(key), Delete(key)
- **Also Fetch based on document contents**

Example systems

- MongoDB, CouchDB, SimpleDB, ...



Graph Database Systems

- Data model: nodes **and** edges
- Nodes may have properties (**including** ID)
- Edges may have labels **or** roles



Graph Database Systems

- Interfaces and query languages vary
- Single-step versus “path expressions” versus full recursion
- Example systems
Neo4j, Pregel, ...
- RDF “triple stores” can map to graph databases



NoSQL Systems

- “NoSQL” = “Not Only SQL”
Not every data management/analysis problem is best solved *exclusively* using a traditional DBMS
- Current incarnations
 - MapReduce framework
 - Key-value stores
 - Document stores
 - Graph database systems

Not Only SQL