

# Data Structures & Algorithms

Adil M. Khan  
Professor of Computer Science  
Innopolis University

# Balanced Binary Search Trees

# Outline



- Motivation
- AVL Trees
- Red-Black Trees

# Binary Search Tree

- For a binary search tree with  $n$  nodes
  - average search and insertion time is  $O(\log n)$
  - It will take  $\log_2 n$  comparisons to find a particular node or find out that it isn't
- However, this is only true if the tree is “balanced”
- That is, the “height” of the tree is balanced

# Binary Search Tree


- In the worst case, insertion and searching time becomes  $O(n)$
- Because the height is  $O(n)$



# Binary Search Tree

- In a dynamic tree, nodes are inserted and deleted over time
- So we must find a way to keep the height of a binary search tree always  $O(\log n)$
- To achieve this, the tree must always be balanced

“for any node, its left subtree should not be much higher than its right subtree, and vice-versa”



# AVL Trees


# AVL Trees

- Adelson-Velskii and Landis in 1962 introduced a binary tree structure that is balanced with respect to the heights of its subtree
- Insertions and deletions are made such that the tree always remain height-balanced



# Height of a Node

- Height of a tree is the maximum over all node depths, or, equivalently, the longest path in the tree
- The height of a tree node  $v$  is defined as the height of the subtree rooted at node  $v$



# AVL Trees

- Definition
- An empty tree is height-balanced
- If **T** is non-empty binary tree with left and right subtrees **T\_1** and **T\_2**

**T** is balanced if and only if

- **T\_1** and **T\_2** are balanced, and
- $|\text{height}(\mathbf{T\_1}) - \text{height}(\mathbf{T\_2})| \leq 1$


# Recall: Binary Tree Terminology

- Height of a tree  $T$
- 0 (for convenience) if  $T$  is empty and  
 $1 + \max(\text{height}(T_1), \text{height}(T_2))$  otherwise,  
where  $T_1$  and  $T_2$  are the subtrees of the root


# Recall: Binary Tree Terminology

- Height Numbering
  - Number all **external (leaf)** nodes **0**
  - Number each **internal** node to be **one more** than the maximum of the numbers of its children
  - Then number of the root node is the height of **T**
- The height of a node **u** in **T** is the height of the subtree rooted at **u**


# AVL Trees




# AVL Trees



# AVL Trees



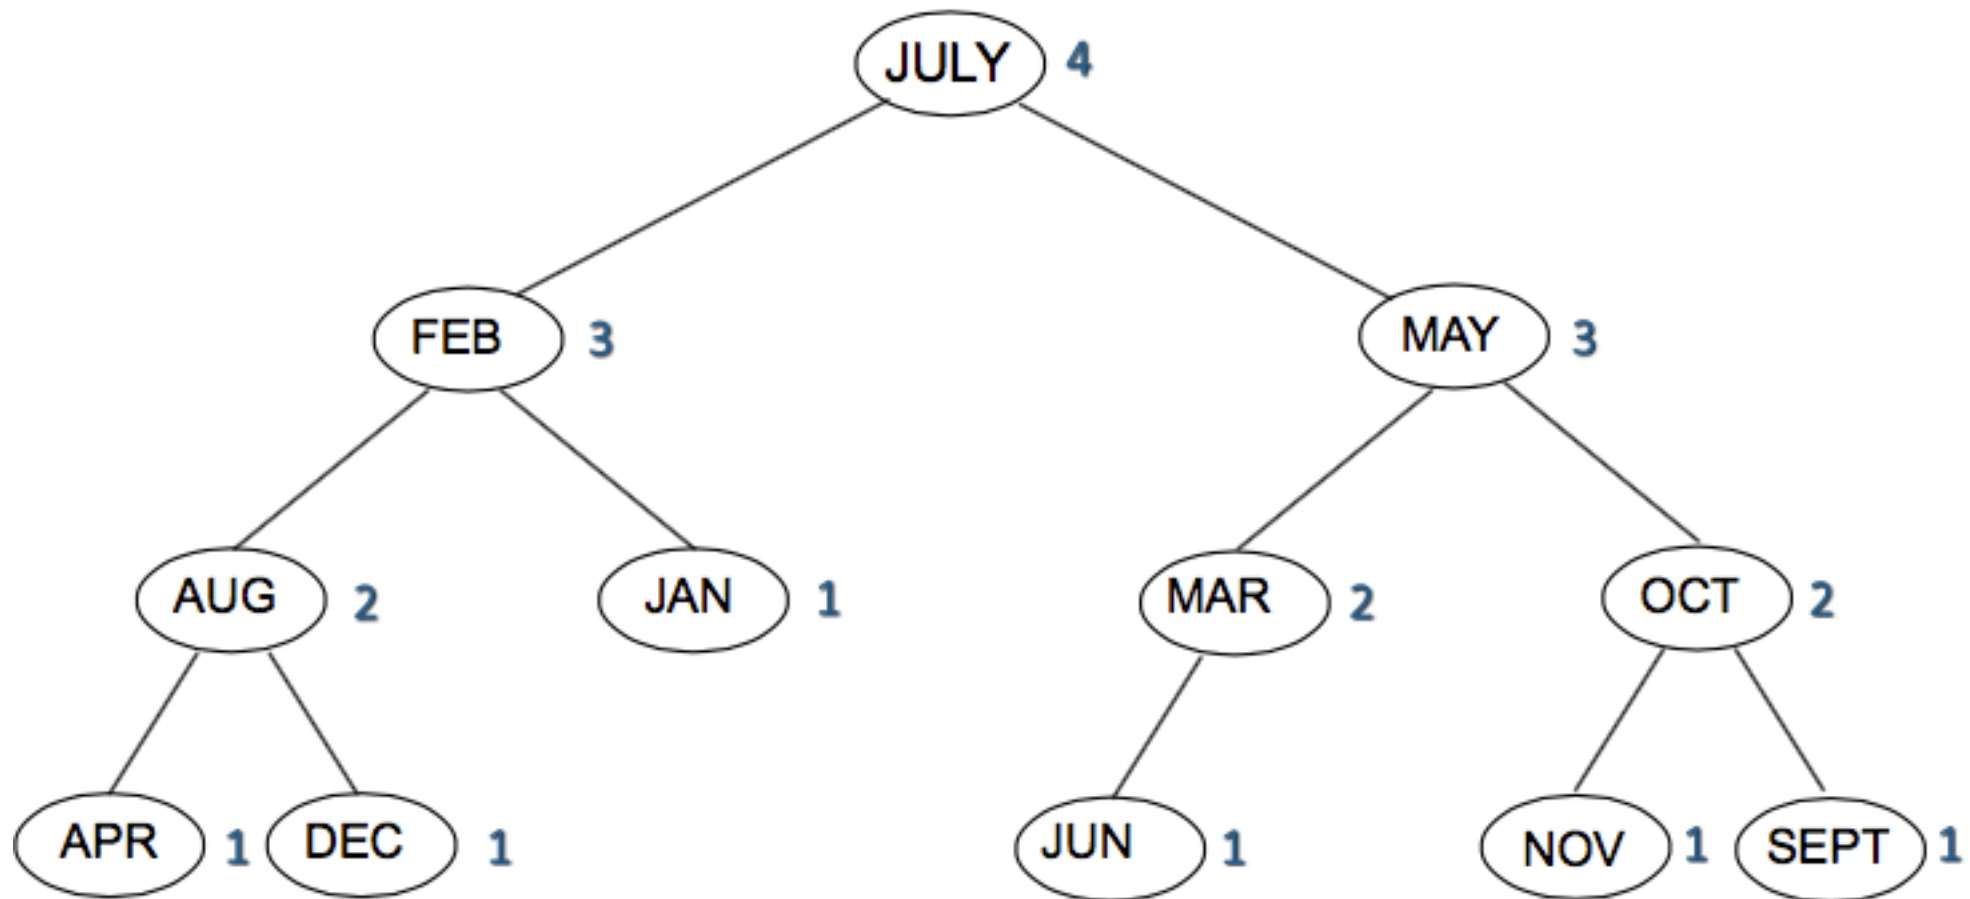
# AVL Trees



A Balanced Tree for the Months of the Year



# AVL Trees




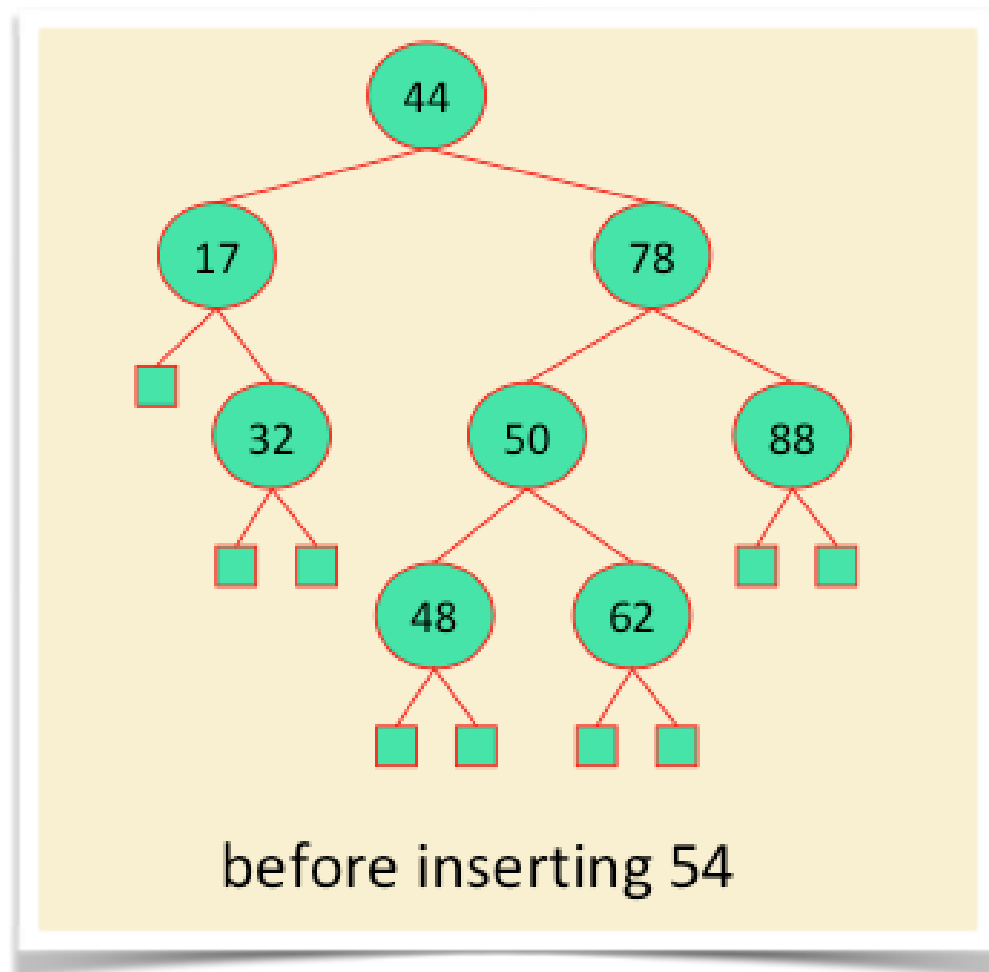
A Balanced Tree for the Months of the Year

# Operations in an AVL Tree

- The height of an AVL tree is  $O(\log n)$
- Thus the **search** operation takes  $O(\log n)$ 
  - Performed just like in a binary search tree since AVL tree is a binary search tree
- What we need to show is how to **insert** and **remove** in AVL trees while maintaining
  - the height balanced property
  - the binary search tree order


# Insertion in an AVL Tree

- Starts as in a binary search tree
- Always done by expanding an external node




# Insertion in an AVL Tree

- After inserting a new node into an AVL tree, the height-balanced property of the AVL tree is very likely lost




# Insertion in an AVL Tree

- Thus, to make it an AVL tree again, we need to **restore the balance** by **restructuring the tree**




# Pictorial Notation



# Restructuring

- Let  $w$  be the new node, just inserted into an AVL tree




# Restructuring

- The next step is to search for the unbalanced node(s)
- Check each node in the tree to see if it is balanced.
- Do you think this approach is efficient?
- Can we make it efficient?




# Restructuring

- Tip: after insertion of **w**, heights could change (increase) only for the ancestors of **w**
- Thus only ancestors of **w** could be unbalanced
- Search up the tree from **w** checking and correcting any unbalanced node




# Restructuring

- Follow the path from **w** to the root





# Restructuring

- Suppose the first unbalanced node is at position **z**
- Height difference between the left and the right subtree of **z** is more than 1
- In fact, it is **exactly 2**
  - tree was balanced before insertion
  - each insertion can change height only by a factor of 1
- **w** is in the higher subtree



# Restructuring

- Two cases:
  - Right subtree is higher
  - Or, left subtree is higher




# Restructuring

- Let **S** be the higher subtree, with height  $p+2$
- Let **y** be the root of **S**


## Refresher

- **w** is the new node
- **z** (ancestor of **w**) is the first unbalanced node
- **S** is the higher subtree of **z**
- **y** is the root of **S**




# Restructuring

- Remember, the tree was balanced before insertion
- So, the height of **S** was  $p+1$  **before insertion**




# Restructuring

- **y** is balanced after insertion, and **z** is not
- So, both subtrees of **S** had height exactly **p** before insertion



# Restructuring


- Complete Picture






# Restructuring

- More Complete Picture :)




# Restructuring

- Let's consider case 1
- Let  $R$  be the right subtree of  $y$
- Let  $x$  be the root of  $R$



# Restructuring




# Restructuring


- R contains w, which is an internal node
- Therefore, R has at least one internal node
- There are two cases:
  - Case A:  $x = w$  in which case  $p = 0$  and both subtrees of w are leaf nodes

# Restructuring

- Case B:
  - height of R went from  $p$  (before insertion) to  $p+1$  after insertion
  - $x$  was balanced before insertion and is balanced after insertion
  - both subtrees of  $x$  had height  $p-1$  before insertion
  - After insertion, one subtree of  $x$  has height  $p$ , the other  $p-1$




# Restructuring




# Restructuring

- Even More Complete Picture :)



# Restructuring


- Finally, let's restructure the tree
- Case 1:






# Restructuring

- Finally, let's restructure the tree
- Case 1:



What's the height differences at nodes **x**, **y** and **z** after restructuring?

# Restructuring




# Restructuring

- All four cases can be coded with the same algorithm called: **Trinode Restructuring**
- Trinode because there are three nodes

# Trinode Restructuring

- In all four cases, out of the three nodes  $x$ ,  $y$ , and  $z$ , make:
  - node with the middle key the new parent
  - smallest key node its left child
  - largest key node its right child
- for the new parent, the previous subtree (if present) must be put in appropriate positions
  - Left subtree (if present) goes with the new left child
  - Right subtree (if present) goes with the new right child



# Trinode Restructuring

- Takes  $O(\log n) + O(1)$
- No loops, no recursive calls, constant number of comparisons, and changes in parent-child relationships
- Only 1 trinode restructuring is needed per insertion to restore the height balance property

# Deletion

- Deletion from an AVL tree may violate the height-balance property, too
- In this case, procedure for restructuring the tree to restore the balance is the same as in the case of insertion, with some changes
  - how to choose  $x$ ,  $y$ , and  $z$
  - repeated restructuring might be needed, max  $O(\log n)$
- For further details, please read section 11.3.1 of your textbooks


# Analysis

- Ok, so we have learned how to keep a binary search tree always balanced (AVL tree) after insertions and deletions
- But why is this important?
- Recall that all we wanted was a way to make and keep the height of a tree with  $n$  nodes  $O(\log n)$
- Is the height of an AVL tree  $O(\log n)$ ?

# Proof

- Proposition: The height of an AVL tree storing  $n$  entries is  $O(\log n)$

Let  $T$  be an AVL tree of height  $h$ .  $T$  can be visualized as






# Proof

Let  $n(h)$  be the minimum number of **internal nodes** in an AVL tree of height  $h$

we know,  $n(1) = 1$  and  $n(2) = 2$

For  $h \geq 3$

$$n(h) = 1 + n(h-1) + n(h-2)$$



# Proof

$$n(h) = 1 + n(h-1) + n(h-2)$$

Now that we know this, the rest is just algebra

According to the properties of Fibonacci progressions

$$n(h) > n(h-1), \text{ so } n(h-1) > n(h-2)$$

By replacing  $n(h-1)$  with  $n(h-2)$  and dropping the 1, we get

$$\mathbf{n(h) > 2\ n(h-2)}$$

# Proof

$$n(h) > 2 n(h-2)$$

We can stop at this point. We have shown that  $n(h)$  at least **doubles** when  $h$  goes **up by 2**. This says that  $n(h)$  is **exponential** in  $h$ , and hence  $h$  is **logarithmic** in  $n$

But let's continue

$$n(h) > 2 (2 n(h-4)) = 2^2 n(h-4)$$

# Proof

Thus,

For any  $i > 0$ ,  $n(h) > 2^i n(h-2i)$

Let's **get rid of  $i$  by expressing it in terms of  $h$** , but choose a value that results in making  $h-2i$  either 1 or 2. It is because we know the values for  $n(1)$  and  $n(2)$

That is, let

$$i = h/2 - 1$$

# Proof

$$n(h) > 2^{h/2-1}n(h-2i) = 2^{h/2-1}$$

By taking logarithmic of both sides

$$\log(n(h)) > (h/2)-1 \text{ or}$$

$$h < 2\log(n(h))+2$$

Or

$$h < \log(n)$$